

Methods	Year	Parameters	Dice	iou	Input size
TransNetR (4)	2023	27.27	87.06	80.16	?
SegMed-Net (8)	2022	59.9	90.8	83.70	352*352
PVT-Cascade (3)	2023	35.2	91.1	86.3	352*352
MegaNet (2)	2024	44.1	91.3	86.3	352*352
Fu-TransHNet (5)	2023	?	91.6	86.6	352*352
LiteBoundD (1)	2026	32.8	91.9	84.1	352*352
AAPCNet (7)	2025	38.2	92.10	86.9	256*256
DVFIP-Net	2026	26.54	02.29	04.78	352*352
SAM2-FNet	2026	2.72	04.29	87.40	352*352
MLB-Net	2026	87.7	92.60	87.80	352*352
MSPN	2026	28.56	92.70	87.90	352*352
HCA-Net	2026	89.70	93.60	89.7	352*352
Ours	2026	38.5	93.89	88.66	352*352
PolyMamba-Net (6)	2026	25.3	94.2	89.1	352*352

1. Sharpening Lightweight Models for Generalized Polyp Segmentation: A Boundary Guided Distillation from Foundation Models
2. “Meganet: Multi-scale edge-guided attention network for weak boundary polyp segmentation
3. “Medical image segmentation via cascaded attention decoding
4. TransNetR: Transformer-based Residual Network for Polyp Segmentation with Multi-Center Out-of-Distribution Testin
5. Cooperation Learning Enhanced Colonic Polyp Segmentation Based on TransformerCNN Fusion
6. PolyMamba-Net: a lightweight and boundary-aware network for real-time polyp segmentation in colonoscopy
7. Attention-Guided Asymmetric Multiscale Polyp Segmentation Network
8. SegMed-Net: a domain-aware dual encoder-decoder network with dense feature fusion for enhanced medical image segmentation

Boundary-Frequency Injection Module:

- Decomposing frequency components
 - Frequency Feature Decomposer (FFD)
 - Spilt skip features into high-frequency and low-frequency using 2D FFT
 - * High-frequency -> Boundaries, edges
 - * Low-frequency -> semantic content
- Using boundary information for better feature guidance
 - Boundary-guided Squeeze-Excitation Injection (BSEI)
 - Uses Sobel operators to detect boundaries in skip features
 - Creates spatial attention maps from boundary information
 - Uses cross-attention between decoder features and boundary-weighted encoder features
- Adaptively fusing these different feature representations
 - Multi-scale Frequency Adaptive Fusion (MFAF)
 - Learns soft weights to combine high-frequency, low-frequency, and boundary-enhanced features

- Uses channel-wise gating to decide the importance of each component

Purpose:

The BFIM replaces standard skip connections to:

- **Better preserve boundary information** during up-sampling
- **Adaptively emphasize** important frequency components
- **Improve segmentation accuracy** especially for object boundaries
- **Provide learnable fusion** instead of simple concatenation

Models	Seen				Unseen						Epoch
	Kvasir-Seg		ClinicDB		CVC-300		ColonDB		ETIS		
	Dice	iou	Dice	iou	Dice	iou	Dice	iou	Dice	iou	
HCA-Net	93.6	89.7	95.6	90.9	91.2	87.3	83.4	76.9	83.3	78.9	
Ours	92.40	86.08	95.19	90.84	91.49	84.32	81.26	70.44	87.39	77.61	1370
Ours-2	92.58	86.41	95.31	91.04			81.45	70.95			1405
	92.03		95.27	90.97			83.34	73.75			

Models	Seen				Unseen						Epoch
	Kvasir-Seg		ClinicDB		CVC-300		ColonDB		ETIS		
	Dice	iou	Dice	iou	Dice	iou	Dice	iou	Dice	iou	
HCA-Net	93.6	89.7	95.6	90.9	91.2	87.3	83.4	76.9	83.3	78.9	
	93.49	87.77	95.31	91.04	91.61	84.52	83.34	73.75	88.69	79.67	
Our	93.49	87.77	94.92	90.33	90.33	82.52	79.72	69.16	88.69	79.67	100
Ours	92.58	86.41	95.31	91.04	91.49	84.32	83.34	73.75	87.39	77.61	1370
Ours-93.04	03.29				52.19	19.38			84.96	73.85	
Ours-93.66	92.75	86.48	95.20	90.83	91.61	84.52	82.81	72.96	85.39	74.50	2836

Introduction

Colorectal polyps are abnormal tissue growths that may develop into malignant lesions if they are not detected and removed at an early stage. Colonoscopy is routinely employed for examination of the lower gastrointestinal tract and enables clinicians to identify and resect suspicious polyps. However, the reliability of visual inspection is affected by operator experience, fatigue, imaging conditions, and the heterogeneous appearance of colorectal lesions. Computer-aided polyp segmentation can assist clinical examination by providing pixel-level localization of suspected lesions and by presenting their spatial extent and boundaries more explicitly.

Accurate polyp segmentation remains challenging because colorectal lesions exhibit substantial variations in size, morphology, texture, color, and location. Some polyps occupy a large and visually distinctive region, whereas others are small, flat, or partially concealed by intestinal folds. In addition, visual similarity between lesion tissue and the surrounding mucosa frequently produces weak or ambiguous boundaries. Specular reflections, motion blur, bubbles, shadows, fecal material, and medical instruments may further interfere with feature extraction. These factors make it difficult for a segmentation model to obtain reliable semantic localization and precise boundary delineation simultaneously.

U-shaped convolutional architectures have been extensively used in medical image segmentation because their contracting paths extract hierarchical semantic features, while their expanding paths progressively recover spatial resolution. Skip connections transfer relatively high-resolution encoder representations to the decoder and therefore facilitate lesion localization. Nevertheless, conventional skip connections may also introduce redundant texture responses and background interference because encoder features are transferred without explicit feature selection. Moreover, repeated spatial down-sampling may remove subtle structural information required for segmenting small or weakly bounded polyps. U-Net established the fundamental contracting-expanding design, while subsequent architectures, including UNet++, redesigned skip pathways and incorporated deep supervision to reduce the semantic discrepancy between encoder and decoder representations [1,2].

To improve global-context modeling, transformer-based and hybrid convolution-transformer architectures have been investigated for medical image segmentation. TransUNet combines convolutional feature extraction with transformer encoding to capture long-range dependencies, while Polyp-PVT employs hierarchical transformer representations and cross-level feature aggregation for polyp segmentation [3,4]. These methods demonstrate the value of global contextual information; however, transformer-based designs can introduce considerable computational complexity and may still require dedicated mechanisms for recovering local boundary information.

Modern convolutional architectures provide an alternative means of obtaining strong hierarchical representations without relying on self-attention. ConvNeXt modernizes conventional convolutional design through large-kernel depthwise convolution, inverted channel expansion, Layer Normalization, GELU activation, layer scaling, and stochastic depth. These characteristics enable ConvNeXt to retain the spatial inductive biases and computational regularity of convolutional networks while providing competitive representation capacity [5].

Despite these advances, three interconnected limitations remain important for polyp segmentation. First, the large variation in lesion size requires contextual representations with multiple receptive fields. Second, weak boundaries require a mechanism that selectively enhances local structural variations rather than indiscriminately transferring skip features. Third, deep semantic representations must be combined with shallow spatial information without allowing noisy low-level responses to dominate the final prediction. Addressing these limitations within a computationally moderate architecture is particularly important for applications in which both accuracy and practical inference efficiency are required.

This study proposes a multi-scale and boundary-sensitive ConvNeXt-UNet for automatic colorectal polyp segmentation. The model employs a ConvNeXt-Tiny encoder to generate four levels of hierarchical representations. A lightweight residual bottleneck and a Multi-Scale Context module are applied to the deepest encoder feature to aggregate local, enlarged-receptive-field, and global information. During decoding, a Boundary-Sensitive Enhancement and Integration module refines the concatenated encoder and decoder features by learning spatial edge-response attention. In parallel, a shallow Detail Branch preserves high-resolution appearance information directly from the input image. A Gated Detail Fusion module selectively injects this detailed representation into the final decoder stage. Three auxiliary output heads provide intermediate supervision during training. The complete architecture contains approximately 29.61 million trainable parameters.

The principal contributions of this work are summarized as follows:

1. A ConvNeXt-Tiny-based U-shaped segmentation architecture is developed to combine hierarchical convolutional feature extraction with a lightweight decoder for colorectal polyp segmentation.
2. A Multi-Scale Context module is introduced at the bottleneck to aggregate features obtained using three dilated depthwise convolution branches and a global-context branch.
3. A Boundary-Sensitive Enhancement and Integration module is designed to refine skip-connected features through learned spatial edge-response modulation.
4. A high-resolution Detail Branch and Gated Detail Fusion module are introduced to preserve shallow spatial information and selectively combine it with the final decoder representation.

5. Intermediate prediction heads are incorporated to support deep supervision and improve optimization of multi-level decoder features.
6. The proposed model is evaluated on Kvasir-SEG, CVC-ClinicDB, CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB using in-distribution and cross-dataset testing protocols.

Related work

Convolutional encoder-decoder architectures

Fully convolutional encoder-decoder networks form the foundation of modern medical image segmentation. U-Net introduced a symmetric contracting and expanding architecture in which high-resolution encoder representations are transferred to the decoder through skip connections. This design enables the network to combine contextual representations with spatially detailed features and has become a standard baseline for biomedical segmentation [1].

Subsequent studies modified the U-shaped architecture to improve feature fusion and optimization. UNet++ introduced nested dense skip pathways to reduce the semantic discrepancy between encoder and decoder features and employed deep supervision at multiple decoder depths [2]. Such developments demonstrated that the design of skip connections is as important as the choice of encoder because direct feature transfer may combine representations with substantially different semantic characteristics.

In polyp segmentation, computationally efficient convolutional networks have also received considerable attention. HarDNet-MSEG combines a low-memory-traffic convolutional backbone with a lightweight decoder and illustrates that competitive segmentation performance can be obtained without a computationally expensive transformer architecture [6]. However, standard convolutional models generally require additional mechanisms to capture broad contextual relationships and refine weak lesion boundaries.

PraNet introduced a parallel partial decoder to aggregate high-level features and a reverse-attention strategy to progressively refine uncertain polyp regions [7]. Its recurrent interaction between regional predictions and boundary cues demonstrated the effectiveness of explicitly directing the decoder toward inadequately segmented regions. Nevertheless, reverse-attention refinement depends on preceding prediction maps and does not directly solve the problem of noisy information transferred through conventional skip connections.

Global and multi-scale contextual modeling

The considerable variation in polyp size motivates the use of multi-scale contextual modeling. A segmentation network must preserve local details for small lesions while maintaining a sufficiently large receptive field to identify extensive or irregularly shaped polyps. Dilated convolutions, pyramid pooling, multi-branch receptive fields, and cross-scale feature aggregation have consequently been incorporated into many segmentation architectures.

Transformer-based approaches address this problem by modeling long-range spatial relationships. TransUNet uses a transformer encoder over convolutional feature representations to capture global context while retaining high-resolution CNN features for localization [3]. Polyp-PVT employs a pyramid vision transformer and specialized fusion modules to integrate semantic and positional information across hierarchical feature levels [4]. These methods improve long-range feature interaction, but their attention operations and cross-level fusion mechanisms may increase architectural and computational complexity.

Recent contextual-information-flow and multi-stage architectures similarly emphasize the need to coordinate local lesion appearance with broader anatomical context. State-space and frequency-oriented architectures have also been investigated. Polyp-Mamba combines multi-frequency perception with a gated selection mechanism to model spatial dependencies and frequency-sensitive representations [8]. The present work does not employ frequency decomposition or a state-space operator; instead, it uses parallel dilated depthwise convolutions and global average pooling to obtain multi-scale contextual information within a purely convolutional framework.

Boundary-aware polyp segmentation

Weak and visually ambiguous lesion boundaries are among the principal difficulties in colorectal polyp segmentation. Boundary-aware networks attempt to preserve or reconstruct structural information that may disappear during repeated down-sampling. These methods commonly introduce explicit edge supervision, auxiliary contour branches, gradient-derived features, reverse attention, or boundary-sensitive feature fusion.

MEGANet employs multi-scale edge-guided attention and incorporates high-frequency information derived from image pyramids to improve segmentation of weak boundaries [9]. Its findings demonstrate that explicitly retaining edge-related information can reduce deterioration of boundary features in deep networks. MSBP-Net performs multi-scale boundary prediction to improve contour delineation while maintaining relatively low model complexity [10]. This line of research indicates that boundary information should be modeled at more than one scale because polyp contours may contain both broad structural transitions and small irregularities.

Explicit edge extraction may also respond to irrelevant structures such as folds, reflections, bubbles, and instruments. Furthermore, additional boundary labels or edge-generation procedures may increase the complexity of the training pipeline. The BSEI module in the proposed model follows a different strategy. It derives a learned spatial response from transformed skip features and uses this response to modulate the features through residual enhancement. Therefore, the method is boundary-sensitive but does not depend on Sobel filtering, Laplacian pyramids, explicit contour labels, or a separate boundary loss.

Detail preservation and adaptive feature fusion

Deep semantic features provide strong category discrimination but have reduced spatial resolution. Conversely, shallow features preserve local appearance and edge information but may contain substantial background noise. Effective segmentation therefore requires controlled interaction between semantic and detailed representations.

Conventional U-shaped networks concatenate or add shallow and deep features directly. Although these operations are computationally simple, they assign no explicit relevance to detailed information. Attention and gating mechanisms provide a more selective alternative by estimating the contribution of one representation conditioned on another. Recent polyp-segmentation architectures employ specialized cross-scale or attention-based fusion to suppress irrelevant responses and improve lesion localization.

The proposed method addresses detail preservation using a shallow Detail Branch operating at one-half of the input resolution. Instead of directly adding this feature to the decoder, the Gated Detail Fusion module estimates a gate from the concatenated decoder and detail representations. The gate modulates the detail feature before final fusion, enabling the network to retain high-resolution information while limiting transfer of irrelevant low-level responses.

Positioning of the proposed method

Existing approaches have improved polyp segmentation through reverse attention, transformer-based contextual modeling, frequency processing, explicit boundary prediction, or complex multi-stage refinement. Nevertheless, many methods address context, boundaries, and fine details as separate problems. The proposed architecture integrates these requirements within a unified convolutional encoder-decoder framework.

Unlike transformer-based methods, the model employs a ConvNeXt-Tiny encoder and depthwise separable convolutions. Unlike frequency-oriented networks, it does not require Fourier or wavelet decomposition. Unlike explicitly supervised boundary networks, it derives a learned local-response map from intermediate features. Finally, unlike direct shallow-deep fusion, it uses a conditioned gate to control the contribution of the high-resolution detail representation. These design choices provide a coherent framework for investigating whether multi-scale context, boundary-sensitive skip refinement, and selective detail fusion can jointly improve polyp segmentation and cross-dataset generalization.

Method

Overall architecture

Let $I \in \mathbb{R}^{3 \times H \times W}$ denote an input colonoscopy image. The proposed network follows a U-shaped encoder-decoder configuration and consists of five principal components: a ConvNeXt-Tiny encoder, a lightweight residual bottleneck, a Multi-Scale Context module, a boundary-sensitive decoder with refined skip connections, and a high-resolution Detail Branch with gated fusion.

The encoder produces four hierarchical feature maps $\{F1, F2, F3, F4\}$ with channel dimensions $\{96, 192, 384, 768\}$ and spatial resolutions $\{H/4 \times W/4, H/8 \times W/8, H/16 \times W/16, H/32 \times W/32\}$. The deepest feature is processed by the lightweight bottleneck and the Multi-Scale Context module. The decoder then progressively increases spatial resolution. At the first three decoding stages, the upsampled decoder feature is concatenated with the corresponding encoder representation and processed by BSEI. At one-half input resolution, the final decoder feature is fused with the output of the Detail Branch. A refinement head converts the fused representation into final segmentation logits, which are bilinearly resized to the original image resolution. During training, three auxiliary prediction heads produce additional full-resolution logits from intermediate decoder stages.

The principal forward operations can be summarized as follows:

$$\begin{aligned} \{F1, F2, F3, F4\} &= E(I), \\ B &= M_MSC(M_LB(F4)), \\ D4 &= M_BSEI,4([U4(B), F3]), \\ D3 &= M_BSEI,3([U3(D4), F2]), \\ D2 &= M_BSEI,2([U2(D3), F1]), \\ D1 &= M_BSEI,1(U1(D2)), \\ P &= U_out(M_ref(M_GDF(D1, M_detail(I)))). \end{aligned}$$

Here, $[\cdot, \cdot]$ denotes channel-wise concatenation, U_i denotes the corresponding decoder up-sampling operation, M_LB is the lightweight bottleneck, M_MSC is the Multi-Scale Context module, M_GDF is Gated Detail Fusion, and P denotes the final segmentation logits.

ConvNeXt-Tiny encoder

Before entering the encoder, the input image is normalized using the ImageNet channel mean and standard deviation:

$$\begin{aligned} \hat{I}_c &= (I_c - \mu_c) / \sigma_c, \\ \mu &= (0.485, 0.456, 0.406), \sigma = (0.229, 0.224, 0.225). \end{aligned}$$

The first encoder stage uses a 4×4 convolution with stride 4, followed by Layer Normalization. The subsequent three stages begin with Layer Normalization and a 2×2 stride-2 convolution. The numbers of ConvNeXt blocks in the four stages are 3, 3, 9, and 3, respectively.

For an input feature X , each ConvNeXt block applies a 7×7 depthwise convolution, converts the representation to channels-last format for Layer Normalization and linear channel projection, expands the channel dimension by a factor of four, applies GELU activation, restores the original channel dimension, and adds the transformed representation through a layer-scaled stochastic residual connection.

$$\begin{aligned} Z &= DWConv_{7 \times 7}(X), \\ V &= W2(GELU(W1(LN(Z)))), \\ Y &= X + DropPath(\gamma \odot V). \end{aligned}$$

Spatial dropout is applied to the outputs of the second, third, and fourth encoder stages. This regularization reduces excessive dependence on individual feature channels in deeper semantic representations.

Lightweight residual bottleneck

The deepest encoder representation F_4 is first processed by a lightweight residual bottleneck. A 1×1 convolution reduces its channel dimension from 768 to 192. A depthwise separable 3×3 convolution performs spatial transformation in the reduced feature space, after which another 1×1 convolution restores the channel dimension to 768.

$$B_0 = F_4 + \gamma_b H(F_4).$$

The transformation $H(\cdot)$ includes pointwise channel reduction, Layer Normalization, GELU activation, depthwise separable convolution, dropout, pointwise channel restoration, and final Layer Normalization. The learnable scalar γ_b is initialized to zero, so the bottleneck initially behaves as an identity mapping and progressively introduces its residual transformation during optimization.

Multi-Scale Context module

The Multi-Scale Context module is applied to B_0 to enrich the bottleneck representation using multiple receptive fields while maintaining moderate computational complexity. A 1×1 convolution first reduces the 768-channel input to 96 channels.

$$R = GELU(LN(Conv1 \times 1(B_0))).$$

Three parallel depthwise 3×3 convolutions with dilation rates $d \in \{1, 3, 5\}$ are applied:

$$C_d = GELU(LN(DWConv_{3 \times 3, d}(R))), \quad d \in \{1, 3, 5\}.$$

A fourth branch extracts image-level context through global average pooling:

$$C_g = Up(GELU(Conv1 \times 1(GAP(R)))).$$

The four branches are concatenated and projected back to 768 channels:

$$C = [C_1, C_3, C_5, C_g],$$

$$B = B_0 + \gamma_c P(C).$$

The projection $P(\cdot)$ consists of a 1×1 convolution, Layer Normalization, GELU activation, and dropout. The scalar γ_c is initialized to zero. Dilation rates 1, 3, and 5 provide local and progressively broader neighborhoods, while the pooled branch supplies image-level context.

Boundary-Sensitive Enhancement and Integration module

At each skip-connected decoder stage, the upsampled decoder feature is concatenated with the corresponding encoder feature. The resulting representation is refined using BSEI. Let $Sl = [Ul, Fl]$ denote the concatenated decoder and encoder features at level l .

$$Ql = GELU(LN(Conv1 \times 1(Sl))),$$

$$Zl = GELU(LN(SepConv3 \times 3(Ql))).$$

A learned depthwise 3×3 convolution generates a local response map. The absolute value represents the magnitude of the learned spatial variation, and a sigmoid bounds the response:

$$Al = sigmoid(|DWConv3 \times 3(Zl)|).$$

The enhanced feature is calculated by residual modulation:

$$Y_l = Z_l + A_l \odot Z_l.$$

Locations with stronger learned spatial responses therefore receive greater amplification. This operation should be described as learned boundary-sensitive or edge-response enhancement, rather than explicit boundary detection. It does not use Sobel filters, manually generated edge maps, channel squeeze-and-excitation, or boundary annotations.

Progressive decoder

Each decoder block consists of bilinear up-sampling by a factor of two, a depthwise separable 3×3 convolution, Layer Normalization, GELU activation, and spatial dropout. The first decoder stage transforms the context-enhanced bottleneck from 768 to 384 channels. After concatenation with F3, BSEI produces a 384-channel representation.

$$D_4 = M_BSEI,4([M_dec,4(B), F3]),$$

$$D_3 = M_BSEI,3([M_dec,3(D_4), F2]),$$

$$D_2 = M_BSEI,2([M_dec,2(D_3), F1]),$$

$$D_1 = M_BSEI,1(M_dec,1(D_2)).$$

The output channel dimensions of D3 and D2 are 192 and 96, respectively. The final decoder block increases the spatial resolution from $H/4 \times W/4$ to $H/2 \times W/2$. No encoder skip feature is concatenated at this final stage because the encoder stem begins at one-quarter input resolution.

High-resolution Detail Branch

The Detail Branch processes the original input independently of the deep encoder. Its objective is to retain relatively high-resolution appearance information that may be attenuated by the stride-4 encoder stem. The branch begins with a 3×3 convolution with stride 2, producing a 32-channel feature at one-half input resolution.

$$T_0 = GELU(BN(Conv_{3 \times 3, s=2}(I))),$$

$$T = GELU(BN(SepConv_{3 \times 3}(T_0))).$$

Dropout is applied after both operations. Unlike the main encoder and decoder, this branch uses Batch Normalization.

Gated Detail Fusion

The final decoder feature D1 and the detail representation T have the same spatial resolution. Before fusion, the detail feature is projected using a depthwise separable convolution:

$$\hat{T} = GELU(BN(SepConv_{3 \times 3}(T))).$$

A spatially varying gate is estimated from the concatenated decoder and detail representations:

$$G = \text{sigmoid}(W_{g2}(GELU(W_{g1}([D_1, \hat{T}])))$$

The gate contains 32 channels and modulates the detail representation:

$$T_g = G \odot \hat{T}.$$

The gated detail representation is concatenated with the decoder feature and projected to 128 channels:

$$F_{fuse} = GELU(LN(Conv_{1 \times 1}([D_1, T_g]))).$$

The gate controls the spatial and channel-wise contribution of the shallow detail representation before it enters the final prediction head.

Prediction head and deep supervision

The fused feature is processed by two depthwise separable convolutions. The first reduces the channel dimension from 128 to 96, and the second reduces it from 96 to 48. A 1×1 convolution then produces the segmentation logits.

$$P0 = \text{Conv}1 \times 1(R(\text{Fuse})),$$

$$P = \text{Up}H, W(P0).$$

For binary segmentation, the probability map at inference is obtained as $\hat{Y} = \text{sigmoid}(P)$. The implementation returns logits rather than sigmoid probabilities; consequently, the activation should be applied by the loss function or by the inference procedure.

During training, auxiliary 1×1 prediction heads are attached to D2, D3, and D4. Their logits are resized to the original input resolution:

$$P2 = \text{Up}H, W(\text{Conv}1 \times 1(D2)),$$

$$P3 = \text{Up}H, W(\text{Conv}1 \times 1(D3)),$$

$$P4 = \text{Up}H, W(\text{Conv}1 \times 1(D4)).$$

A general deeply supervised objective can be written as:

$$L_{total} = L_{seg}(P, Y) + \lambda_2 L_{seg}(P2, Y) + \lambda_3 L_{seg}(P3, Y) + \lambda_4 L_{seg}(P4, Y).$$

Here, Y denotes the reference mask and λ_2 , λ_3 , and λ_4 control the auxiliary losses. The exact segmentation loss and numerical auxiliary weights are not defined in the supplied architecture file and must be obtained from the training implementation before the manuscript is finalized.

Model complexity

For a single-class output configuration, the supplied implementation contains 29,609,302 trainable parameters (approximately 29.61 million). This value is consistent with the approximately 29.60 million parameters reported in the later complexity table of the experimental draft. Any earlier table reporting 38.5 million parameters should be corrected or explicitly identified as belonging to a different model version.

Experiments

Dataset:

To verify the effectiveness of the models, we adopted five public datasets (including Kvasir-SEG [ho-32], CVC-ClinicDB [ho-33], ETIS [ho-34], CVC-ColonDB [ho-35] and CVC-300 [ho-36]) commonly used in existing polyp segmentation studies. These images are acquired from colonoscopy video sequences, and all of them contain the corresponding ground truth. For fair comparison, according to table 1, we split the datasets as in PraNet [8]. The training set consists of 900 images from the Kvasir-SEG and 550 images from the CVC-ClinicDB. The remaining samples in the Kvasir-SEG and the CVC-ClinicDB, and the other three datasets are used as five independent test sets. In this study, we randomly selected 145 images (10 % of the training set) from the training set as the validation set, serving as the criterion for saving the model during training.

Table 1 Number, resolution, and trainin/test split of the five public polyp segmentation datasets. Resized to 352×352 .

Dataset	Number	Train	Test	Resolution
Kvasir-Seg	1000	900	500	500*574
CVC-ClinicDB	612	550	62	288*384
CVC-ColonDB	380	-	380	500*570

Dataset	Number	Train	Test	Resolution
ETIS-LaribPolypDB	196	-	196	966*1225
CVC-300	60	-	60	500*570

Loss function

Evaluation metrics

To objectively evaluate model performance, two commonly used segmentation metrics are employed: the mean Intersection over Union (mIoU) and mean Dice coefficient (mDice). In addition, in order to evaluate the performance of the model more comprehensively, this paper also adds four evaluation indicators commonly used in the field of salient object detection, namely S-measure (S_α) [ho-f20], F-measure [], E-measure () [ho-fan] and mean absolute error (MAE) [ho-peraz]. The detailed definitions and formulas for these metrics are as follows: mIoU and mDice are commonly used to measure the similarity between the prediction map and the real map. They quantitatively analyze the spatial overlap between the prediction map and the real map. The higher the value, the better the segmentation accuracy of the model. Among them, the average means that the IoU coefficient and Dice coefficient of each category are averaged, but there is only one category of polyps here. Thus, their calculation formulas are as follows:

$$mIoU = \frac{TP}{TP + FN + FP}$$

$$mDice = \frac{2 * TP}{2 * TP + FP + FN}$$

Where TP (True Positive) indicates the real sample, that is, the forecast is a positive sample, and the actual sample is also a positive sample; FP (False Positive) indicates the false positive sample, that is, the prediction is a positive sample, and the actual sample is a negative sample; FN (False Negative) indicates the false negative sample, that is, the prediction is a negative sample, and the actual sample is a positive sample.

S-measure, S_α , is used to measure the structural similarity between region perception (S_r) and object perception (S_o) between the predicted graph and the real graph. It comprehensively considers the similarity of spatial structure and regional contour, and is a comprehensive evaluation of segmentation performance. Thus, this calculation formulas are as follows:

$$S_\alpha = \alpha * S_o + (1 - \alpha) * S_r$$

Where α is the weight parameter, and the default setting is 0.5.

F-measure, F_β , comprehensively considers the importance of precision and recall rate, and gives corresponding weights to precision and recall rate according to specific standards. It provides a balanced evaluation index for the evaluation model. Thus, F_β calculation formula is as follows:

$$F_\beta = (1 + \beta^2) \frac{P * R}{\beta^2 P + R}$$

Where P and R are precision and recall respectively; β^2 parameter, which is set to 0.3.

E-measure, E_φ , is used to quantify alignment between predicted and ground truth masks at image and pixel levels. Its calculation formula is as follows:

$$E_\varphi = \frac{1}{W * H} \sum_{i=1}^W \sum_{j=1}^H \phi_{FM}(i, j),$$

Where ϕ_{FM} is the enhanced alignment matrix [ho-40].

MAE is used to compute the mean absolute error between predicted and ground truth masks. The calculation formula is as follows:

$$MAE = \frac{1}{W * H} \sum_{i=1}^W \sum_{j=1}^H |T(i, j) - G(i, j)|,$$

Where T and G denote the predicted and ground truth masks, respectively.

Implementation details

In this work, the proposed ConvNeXt-UNet was trained on a GeForce RTX 3090 GPU with 24 GB memory and a pytorch framework. the AdamW optimizer was employed with a learning rate of 0.0001 and weight decay of 0.0004, and the batch size was set to 16. All samples were resized to 352×352 before training. An early stopping mechanism was used to supervise the training. Experiments indicated that the model converged to its optimal state at approximately the 120th epoch.

Comparative study

Quantitative analysis

Qualitative analysis

Models	Seen		ClinicDB		Unseen		ColonDB		ETIS		Epoch
	Kvasir-Seg		Dice	iou	CVC-300		Dice	iou	Dice	iou	
HCA-Net	93.6	89.7	95.6	90.9	91.2	87.3	83.4	76.9	83.3	78.9	
Our9388	93.49	87.77	94.92	90.33	90.33	82.52	79.74	69.20	88.69	79.67	100
Ours	92.58	86.41	95.31	91.04	91.49	84.32	83.34	73.75	87.39	77.61	1370
Ours-93.04	03.29				52.19	19.38			84.96	73.85	
Ours-93.66	92.75	86.48	95.20	90.83	91.61	84.52	82.81	72.96	85.39	74.50	2836

Table 1 Quantitative comparison of various models on the Kvasir-SEG dataset. The best results are highlighted in bold.

Methods	Year	mDice	mIOU	$\mathbf{F}_{\beta}^w$	$\mathbf{S}_{\alpha}$	$\mathbf{mE}_{\xi}$	$\mathbf{maxE}_{\xi}$	MAE
CTNet[20-1]	2024	0.917	0.863	0.910	0.928	0.959	0.964	0.023
MEGANet[47-1]	2024	0.913	0.863	0.907	0.918	0.956	0.959	0.025
CAFÉ-Net[2]	2024	0.933	0.889	0.927	0.939	0.967	?	0.019
Polyp-LVT[2]	2024	0.909	0.851	0.913	0.922	0.941	?	0.024
Polyp-Mamba[48-1]	2025	0.919	0.867	0.912	0.951	-	0.968	0.021
MEIN[25-1]	2025	0.905	0.847	0.905	0.914	-	0.964	0.026
MF-Net[2]	2025	0.915	0.863	0.939	0.924	0.962	?	0.022
MSBP-Net[2]	2025	0.919	0.868	0.916	0.926	0.961	?	0.023
CIFFormer[2]	2025	0.926	0.876	0.929	0.939	0.967	?	0.019
PFPRNet[2]	2025	0.930	0.881	0.920	0.935	0.971	?	0.019

Methods	Year	mDice	mIOU	$\mathbf{F}_\beta^w$	$\mathbf{S}_\alpha$	$\mathbf{mE}_\xi$	$\max \mathbf{E}_\xi$	MAE
PraNet-V2[11-3]	2025	0.907	0.853	?	?	?	?	0.025
MLB-Net[1]	2026	0.926	0.878	0.921	0.938	0.969	0.975	0.021
MSPN[3]	2026	0.927	0.879	?	?	?	?	0.021
DVFIP-Net[4]	2026	0.922	0.874	?	?	?	?	?
Ours	-	0.935	0.879	0.901	0.935	0.952	0.963	0.025

Table 2 Quantitative comparison of various models on the CVC-ClinicDB dataset. The best results are highlighted in bold.

Methods	Year	mDice	mIOU	$\mathbf{F}_\beta^w$	$\mathbf{S}_\alpha$	$\mathbf{mE}_\xi$	$\max \mathbf{E}_\xi$	MAE
CTNet[20-1]	2024	0.936	0.887	0.934	0.952	0.983	0.986	0.006
MEGANet[47-1]	2024	0.938	0.894	0.940	0.950	-	0.987	0.007
Polyp-Mamba[48-1]	2025	0.941	0.896	0.936	0.970	-	0.970	0.008
MEIN[25-1]	2025	0.923	0.878	0.928	0.944	-	0.970	0.008
MF-Net[2]	2025	0.931	0.884	0.953	0.949	0.983	?	0.006
MSBP-Net[2]	2025	0.940	0.892	0.927	0.939	0.976	?	0.007
CIFFormer[2]	2025	0.944	0.897	0.941	0.934	0.972	?	0.007
PFPRNet[2]	2025	0.949	0.903	0.947	0.957	0.989	?	0.006
PraNet-V2[11-3]	2025	0.929	0.88	?	?	?	?	0.009
MLB-Net[1]	2026	0.945	0.902	0.948	0.958	0.989	0.992	0.006
MSPN[3]	2026	0.943	0.902	?	?	?	?	0.006
DVFIP-Net[4]	2026	0.942	0.896	?	?	?	?	?
Ours	-	0.951	0.906	0.931	0.954	0.983	0.990	0.009

Table 3 Quantitative comparison of various models on the ETIS-LaribPolypDB dataset. The best results are highlighted in bold.

Methods	Year	mDice	mIOU	$\mathbf{F}_\beta^w$	$\mathbf{S}_\alpha$	$\mathbf{mE}_\xi$	$\max \mathbf{E}_\xi$	MAE
CTNet[20-1]	2024	0.810	0.734	0.776	0.886	0.913	0.921	0.014
MEGANet[47-1]	2024	0.789	0.709	0.753	0.866	0.855	0.915	0.015
Polyp-Mamba[48-1]	2025	0.756	0.668	0.756	0.887	-	0.872	0.014
MEIN[25-1]	2025	0.788	0.711	0.764	0.870	-	0.902	0.018
MF-Net[2]	2025	0.821	0.750	0.850	0.901	0.925	?	0.011
MSBP-Net[2]	2025	0.795	0.718	0.770	0.873	0.899	?	0.019
CIFFormer[2]	2025	0.810	0.746	0.858	0.889	0.918	?	0.013
PFPRNet[2]	2025	0.828	0.764	0.839	0.908	0.935	?	0.012
PraNet-V2[11-3]	2025	0.713	0.633	?	?	?	?	0.021
MLB-Net[1]	2026	0.821	0.744	0.780	0.899	0.914	0.935	0.013
MSPN[3]	2026	0.798	0.717	?	?	?	?	0.017
DVFIP-Net[4]	2026	0.811	0.730	?	?	?	?	?
Ours	-	0.887	0.797	0.730	0.873	0.883	0.913	0.014

Table 4 Quantitative comparison of various models on the CVC-ColonDB dataset. The best results are highlighted in bold.

Methods	Year	mDice	mIOU	$\mathbf{F}_\beta^w$	$\mathbf{S}_\alpha$	$\mathbf{mE}_\xi$	$\max \mathbf{E}_\xi$	MAE
CTNet[20-1]	2024	0.813	0.734	0.801	0.874	0.915	0.919	0.040
MEGANet[47-1]	2024	0.793	0.714	0.779	0.854	0.892	0.895	0.034
Polyp-Mamba[48-1]	2025	0.791	0.713	0.765	0.846	-	0.902	0.034
MEIN[25-1]	2025	0.791	0.715	0.785	0.856	-	0.897	0.034
MF-Net[2]	2025	0.825	0.751	0.864	0.887	0.929	?	0.024
MSBP-Net[2]	2025	0.810	0.731	0.810	0.861	0.912	?	0.029
CIFFormer[2]	2025	0.823	0.741	0.851	0.874	0.918	?	0.026
PFPRNet[2]	2025	0.819	0.741	0.843	0.882	0.915	?	0.028
PraNet-V2[11-3]	2025	0.702	0.620	?	?	?	?	0.044
MLB-Net[1]	2026	0.832	0.757	0.825	0.886	0.924	0.931	0.023
MSPN[3]	2026	0.811	0.734	?	?	?	?	0.029
DVFIP-Net[4]	2026	0.826	0.746	?	?	?	?	?
Ours	-	79.74	69.20	0.759	0.869	0.885	0.903	0.035

Table 5 Quantitative comparison of various models on the CVC-300(Endoscene) dataset. The best results are highlighted in bold.

Methods	Year	mDice	mIOU	$\mathbf{F}_\beta^w$	$\mathbf{S}_\alpha$	$\mathbf{mE}_\xi$	$\max \mathbf{E}_\xi$	MAE
CTNet[20-1]	2024	0.908	0.844	0.894	0.975	0.975	0.982	0.006
MEGANet[47-1]	2024	0.899	0.834	0.882	0.935	0.966	0.969	0.007
Polyp-Mamba[48-1]	2025	0.906	0.840	0.888	0.936	-	0.975	0.006
MEIN[25-1]	2025	0.893	0.828	0.880	0.930	-	0.965	0.007
MF-Net[2]	2025	0.897	0.836	0.918	0.941	0.966	?	0.007
MSBP-Net[2]	2025	0.903	0.859	0.927	0.942	0.967	?	0.019
CIFFormer[2]	2025	0.883	0.822	0.912	0.939	0.961	?	0.10
PFPRNet[2]	2025	0.891	0.833	0.917	0.928	0.950	?	0.013
PraNet-V2[11-3]	2025	0.873	0.795					0.007
MLB-Net[1]	2026	0.915	0.852	0.909	0.974	0.974	0.978	0.005
MSPN[3]	2026	0.907	0.840					0.006
DVFIP-Net[4]	2026	0.900	0.833	?	?	?	?	?
Ours	-	0.916	0.845	0.848	0.939	0.951	0.969	0.007

Ablation Studies

To systematically evaluate the individual contributions of the proposed modules to the overall segmentation performance, a comprehensive ablation study was conducted. We define our baseline model as the ConvNeXt-Tiny encoder paired with a lightweight decoder and the final prediction head, completely omitting the Multi-Scale Context (MSC) module, BSEI, the Detail Branch (DB), Gated Detail Fusion (GDF), and Deep Supervision (DS). By incrementally adding or systematically removing these components, we isolated their specific impacts on the network. The quantitative results of these architectural modifications are summarized in Table XX.

Effectiveness of MSC

The Multi-Scale Context (MSC) module is designed to capture and aggregate multi-scale contextual information, which is critical for handling significant variations in target object sizes. To assess its impact, we

compared the baseline model with a configuration integrating the MSC module. As shown in Table XX, the addition of the MSC module effectively enhanced the network's spatial adaptability, yielding an improvement of XX% in mDice and XX% in mIoU on the XX dataset. This confirms the module's capability to enrich feature representations prior to the decoding stage.

Effectiveness of BSEI

The BSEI module is introduced to [describe specific function, e.g., enhance boundary spatial features / integrate edge-level information] during the feature extraction process. In our ablation experiments, removing the BSEI module from the combined network led to a noticeable degradation in segmentation accuracy, particularly in [describe specific area, e.g., complex boundary regions]. Specifically, the performance metrics decreased by XX% in mDice and XX% in mIoU (Table XX). This demonstrates that the BSEI module successfully contributes to [describe specific benefit, e.g., refining target boundaries and reducing false positives].

Effectiveness of DB

The Detail Branch (DB) was specifically developed to preserve high-resolution spatial details that are typically lost as features pass through the deeper, heavily down-sampled stages of the encoder. Discarding the DB in the ablation experiments resulted in overly smoothed predictions and a loss of fine-grained structural delineation. Table XX reveals that incorporating the DB leads to an XX% gain in mDice on the XX dataset, validating its necessity for extracting and maintaining high-fidelity spatial details throughout the network.

Effectiveness of GDF

The Gated Detail Fusion (GDF) module acts as a selective mechanism to optimally combine the high-resolution spatial features from the DB with the high-level semantic features from the main network. To isolate the contribution of this gating mechanism, we replaced the GDF module with a standard feature fusion strategy, utilizing simple concatenation and element-wise addition. The results in Table XX indicate that the GDF module significantly outperforms simple fusion, achieving an XX% improvement in mIoU. This highlights the GDF's effectiveness in filtering out background noise and dynamically emphasizing salient target features during the fusion process.

Effectiveness of DS

Finally, Deep Supervision (DS) is employed during the training phase to facilitate direct gradient flow to intermediate layers, thereby alleviating the vanishing gradient problem and optimizing multi-level feature learning. When DS was omitted from the complete proposed architecture (evaluating only the final prediction head loss), the model exhibited slower convergence and a drop in final segmentation accuracy. As evidenced in Table XX, the inclusion of DS improved the mDice by XX%, ensuring more robust feature optimization and contributing to the superior generalization of the proposed framework.

Table X. Ablation study of the proposed architectural components on the five test sets.

Dataset	Metrics Params(M) GFlops(G)	Model Baseline	+MSC	+BSEI	+DB	+GDF	+DS
Kvasir-Seg	mDice $\uparrow$						
	mIoU $\uparrow$						
	$\mathbf{F}_\beta^w \uparrow$						
	$\mathbf{S}_\alpha \uparrow$						
	$\mathbf{mE}_\xi \uparrow$						
	$\max \mathbf{E}_\xi \uparrow$						
	MAE $\downarrow$						

Dataset	Metrics Params(M) GFlops(G)	Model Baseline	+MSC	+BSEI	+DB	+GDF	+DS
CVC- ClinicDB	mDice $\uparrow$						
	mIoU $\uparrow$						
	$\mathbf{F}_{\beta}^w \uparrow$						
	$\mathbf{S}_{\alpha} \uparrow$						
	$\mathbf{mE}_{\xi} \uparrow$						
	$\max \mathbf{E}_{\xi} \uparrow$						
	MAE $\downarrow$						
CVC- ColonDB	mDice $\uparrow$						
	mIoU $\uparrow$						
	$\mathbf{F}_{\beta}^w \uparrow$						
	$\mathbf{S}_{\alpha} \uparrow$						
	$\mathbf{mE}_{\xi} \uparrow$						
	$\max \mathbf{E}_{\xi} \uparrow$						
	MAE $\downarrow$						
CVC-300	mDice $\uparrow$						
	mIoU $\uparrow$						
	$\mathbf{F}_{\beta}^w \uparrow$						
	$\mathbf{S}_{\alpha} \uparrow$						
	$\mathbf{mE}_{\xi} \uparrow$						
	$\max \mathbf{E}_{\xi} \uparrow$						
	MAE $\downarrow$						
ETIS	mDice $\uparrow$						
	mIoU $\uparrow$						
	$\mathbf{F}_{\beta}^w \uparrow$						
	$\mathbf{S}_{\alpha} \uparrow$						
	$\mathbf{mE}_{\xi} \uparrow$						
	$\max \mathbf{E}_{\xi} \uparrow$						
	MAE $\downarrow$						

Notes. MSC = Multi-Scale Context module; DB = Detail Branch; GDF = Gated Detail Fusion; DS = Deep Supervision. Baseline denotes the ConvNeXt-Tiny encoder with the lightweight decoder and final prediction head, without MSC, BSEI, DB, GDF, or DS.

Discussion

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

References

1. MLB-Net
2. HCA-Net

3. MSPN - FROMCOARSETOFINE:MULTI-STAGEPROGRESSIVENETWORKFORCOLONPOLYP SEGMENTATION

- [20-1] Xiao B, HuJ,LiW,PunCM,BiX.Ctnet:Contrastive transformer network for polyp segmentation. IEEE Transactions on Cybernet ics. 2024;.
- [47-1] Bui NT,HoangDH,NguyenQT,TranMT,LeN.Meganet:Multi scale edge-guided attention network for weak boundary polyp segmentation. In: Proceedings of the IEEE/CVF Winter Confer ence on Applications of Computer Vision; 2024. p. 7985–7994.
- [48-1] Zhu X, Wang W, Zhang C, Wang H. Polyp-Mamba: A Hybrid Multi-Frequency Perception Gated Selection Network for polyp segmentation. Information Fusion. 2025;115:102759.
- [25-1] Kang X, MaZ, Liu K, Li Y, Miao Q. Modeling multi-scale uncer tainty with evidence integration for reliable polyp segmentation. Neural Networks. 2025;p. 107553.
- [11-3] Bo-Cheng Hu, Ge-Peng Ji, Dian Shao, and Deng Ping Fan, “Pranet-v2: Dual-supervised reverse atten tion for medical image segmentation,” arXiv preprint arXiv:2504.10986, 2025. 1, 4
- [25-2] S. Wang, S. Lin, F. Sun, and X. Li, “Multi-feature fusion for accurate polyp segmentation using pyramid visual transformers,” Expert Systems with Applications, vol. 280, p. 127558, 2025.
- [34-2] X.-L. Pan, J.-R. Ding, X. Li, S. Liu, J. Wang, B. Hua, G.-Z. Tang, and C.-H. Zhong, “Msbp-net: A multi-scale boundary prediction network for automated polyp segmentation,” Pattern Recognition, vol. 170, p. 112101, 2026.
- [26-2] C. Xu, L. Lin, B. Wang, and J. Liu, “Ciformer: A contextual information flow guided transformer for colorectal polyp segmentation,” Neurocom puting, p. 130413, 2025.
- [29-2] J. Chu, W. Liu, Q. Tian, and W. Lu, “Pfprnet: A phase-wise feature pyramid with retention network for polyp segmentation,” IEEE Journal of Biomedical and Health Informatics, 2024.
- [ho-32] D. Jha, P.H. Smedsrud, M.A. Riegler, et al., Kvasir-SEG: a segmented polyp dataset, in: 26th International conference on MultiMedia Modeling, 2020, pp. 451–462, https://doi.org/10.1007/978-3-030-37734-2_37.
- [ho-33] J. Bernal, F.J. S´anchez, G. Fern´andez-Esparrach, et al., WM-DOVA maps for accurate polyp highlighting in colonoscopy: validation vs. saliency maps from physicians, Comput. Med. Imaging Graph. 43 (2015) 99–111, <https://doi.org/10.1016/j.compmedimag.2015.02.007>.
- [ho-34] J. Silva, A. Histace, O. Romain, et al., Toward embedded detection of polyps in WCE images for early diagnosis of colorectal cancer, Int J Comput Assist Radiol Surg 9 (2) (2014) 283–293, <https://doi.org/10.1007/s11548-013-0926-3>.
- [ho-35] N. Tajbakhsh, S.R. Gurudu, J. Liang, Automated polyp detection in colonoscopy videos using shape and context information, IEEE Trans Med Imaging 35 (2) (2016) 630–644, <https://doi.org/10.1109/TMI.2015.2487997>.
- [ho-36] D. V´azquez, J. Bernal, J.F. S´anchez, et al., A benchmark for endoluminal scene segmentation of colonoscopy images, J Heal. Eng 2017 (2017) 1–9, <https://doi.org/10.1155/2017/4037190>.
- [ho-f20]

Selected References and Citation Placeholders

The entries below support the current draft structure. Complete and standardize all metadata using the journal's required reference style and a reference manager before submission.

- [1] O. Ronneberger, P. Fischer, and T. Brox, ‘U-Net: Convolutional networks for biomedical image segmentation,’ Medical Image Computing and Computer-Assisted Intervention, 2015.
- [2] Z. Zhou, M. M. R. Siddiquee, N. Tajbakhsh, and J. Liang, ‘UNet++: Redesigning skip connections to exploit multiscale features in image segmentation,’ IEEE Transactions on Medical Imaging, 2020.

- [3] J. Chen et al., ‘TransUNet: Transformers make strong encoders for medical image segmentation.’
- [4] B. Dong et al., ‘Polyp-PVT: Polyp segmentation with pyramid vision transformers.’
- [5] Z. Liu et al., ‘A ConvNet for the 2020s,’ Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2022.
- [6] C.-H. Huang et al., ‘HardNet-MSEG: A simple encoder-decoder polyp segmentation neural network that achieves over 0.9 mean Dice and 86 FPS.’
- [7] D.-P. Fan et al., ‘PraNet: Parallel reverse attention network for polyp segmentation,’ Medical Image Computing and Computer-Assisted Intervention, 2020.
- [8] X. Zhu, W. Wang, C. Zhang, and H. Wang, ‘Polyp-Mamba: A hybrid multi-frequency perception gated selection network for polyp segmentation,’ Information Fusion, vol. 115, article 102759, 2025.
- [9] N. T. Bui, D. H. Hoang, Q. T. Nguyen, M. T. Tran, and N. Le, ‘MEGANet: Multi-scale edge-guided attention network for weak boundary polyp segmentation,’ Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision, pp. 7985–7994, 2024.
- [10] X.-L. Pan et al., ‘MSBP-Net: A multi-scale boundary prediction network for automated polyp segmentation,’ Pattern Recognition, vol. 170, article 112101, 2026.

Methods	Year	Param (M)	GFLOPs (G)
CTNet[20-1]	2024		
MEGANet[47-1]	2024		
Polyp-Mamba[48-1]	2025	49.5	27.9
MEIN[25-1]	2025		
MF-Net[25-2]	2025		
MSBP-Net[34-2]	2025	25.52	12.86
CIFFormer[26-2]	2025	32.44	13.14
PFPRNet[29-2]	2025	118.52	39.39
PraNet-V2[11-3]	2025		
MLB-Net[1]	2026	87.70	
MSPN[3]	2026	28.56	13.21
DVFIP-Net[4]	2026	46.62	19.37
Ours	-	29.60	14.57




Fig. 6 Attention heatmaps for the model's MPDF, SSFA_i, and MDI modules. Darker colors indicate higher attention

- **MultiScaleContext**: tests whether multi-scale/global context helps.
- **BSEI**: tests boundary-aware feature enhancement.
- **DetailBranch**: tests high-resolution spatial detail recovery.
- **GatedDetailFusion**: tests whether adaptive fusion is better than simple fusion.




Fig. 1 Six examples of polyp images, with the circled green area in each image representing the polyp




Fig. 8. Module visualization feature map.

TABLE VII
ABLATION STUDY ON BACKBONE WIDTH (BASE_CHANNEL) AND THE COMPUTATIONAL OVERHEAD OF HCA.

base_channel	HCA	CVC-ClinicDB		ETIS-LaribPolypDB		Param (MB)	GFLOPs (G)
		mDice	mIoU	mDice	mIoU		
7	–	0.801	0.748	0.662	0.627	3.45	2.15
7	+	0.902	0.847	0.792	0.742	5.12	3.39
10	–	0.822	0.773	0.693	0.642	6.82	4.05
10	+	0.928	0.879	0.812	0.763	8.94	6.01
12	–	0.843	0.797	0.721	0.672	7.92	6.21
12	+	0.953	0.909	0.833	0.790	10.13	9.33
14	–	0.846	0.802	0.724	0.676	10.77	8.43
14	+	0.956	0.911	0.837	0.794	13.68	12.09
17	–	0.851	0.805	0.729	0.681	15.91	12.11
17	+	0.958	0.913	0.840	0.799	20.36	17.98




Fig. 6. Performance vs. size on the five polyp datasets. The x-axis denotes the proportion (%) of polyp size, and the y-axis represents the averaged mDice coefficient.




Fig. 4 Our model was compared with the results of various advanced methods on the polyp dataset

- Image(Diagram) of the model with related Modules -> two images
- Heatmap of the model in ablation study with 6 samples
- Heatmap of the model for some stages

- Heatmap of the model in comparison with other models
- Table of the ablation study