

Cross-Level Residual Frequency-Guided Multi-Scale Refinement with a ConvNeXt Encoder for Polyp Segmentation

Abstract

Accurate colorectal polyp segmentation remains challenging because lesions vary in scale, contrast, and boundary appearance. We present a ConvNeXt-Tiny U-shaped framework with bottleneck Frequency-Aware Feature Enhancement (FAFEM) and residual frequency-guided coupling to a Multi-Scale Convolution Block (MSCB). The resulting RFG-MSCB uses FAFEM’s low- and high-frequency descriptor to weight multi-kernel branches at the Stage-3 encoder–decoder fusion while preserving ordinary MSCB aggregation as a reference. Under a PraNet-style protocol, checkpoint selection uses only an internal validation subset, with Kvasir-SEG and CVC-ClinicDB subsets reserved for final seen-domain testing and three datasets used for external testing. The model achieved mDice/mIoU of 0.932/0.883 on Kvasir-SEG and 0.948/0.904 on CVC-ClinicDB; external-set mDice ranged from 0.816 to 0.916. Across the five datasets, Residual RFG-MSCB increased aggregate mDice from 0.869 to 0.887 over FAFEM + MSCB while adding 0.08 M parameters, with no additional GFLOPs at the displayed precision.

Introduction

Colorectal polyps are abnormal tissue growths that can develop into malignant lesions if they are not detected and removed at an early stage. Colonoscopy is widely used for the examination of the lower gastrointestinal tract and enables clinicians to identify and remove suspicious lesions. Nevertheless, visual inspection remains challenging because polyp appearance varies considerably across patients and imaging conditions. Variations in lesion size, morphology, texture, contrast, and location, together with specular reflections, motion blur, intestinal folds, bubbles, and other visual artifacts, can complicate accurate delineation of polyp regions. Automated polyp segmentation therefore remains an important problem in computer-aided colonoscopy, where reliable pixel-level localization may assist subsequent analysis and clinical decision support.

Accurate polyp segmentation requires a model to simultaneously capture high-level semantic information and preserve sufficiently detailed spatial representations. U-shaped encoder–decoder architectures are well suited to this task because they combine progressively abstract encoder features with higher-resolution representations through skip connections. U-Net established this general design, while subsequent architectures such as UNet++ further demonstrated the importance of multi-level feature interaction and hierarchical decoding [1, 2]. However, repeated downsampling reduces spatial resolution, and fixed decoder operations may not always provide sufficient flexibility for lesions that vary considerably in scale and visual frequency content.

Recent transformer-based and hybrid architectures have improved long-range dependency modeling in medical image segmentation. TransUNet combines convolutional representations with transformer encoding, while Polyp-PVT applies hierarchical transformer features and cross-level aggregation to polyp segmentation [3, 4]. These approaches demonstrate the importance of broad contextual modeling, but they can introduce additional architectural complexity. ConvNeXt provides a convolutional alternative with hierarchical representations and large-kernel depthwise spatial processing [5].

Multi-scale representation is particularly important in polyp segmentation because lesions may occupy very small regions or extend over a large portion of the image. Multi-Scale Convolution Blocks (MSCBs) address this issue by combining convolutional branches with different receptive fields, allowing a network to aggregate local and broader spatial responses [6]. In their ordinary form, the branch responses are summed without sample-specific frequency conditioning.

Frequency-aware feature processing provides another complementary perspective. We adopt a locally implemented bottleneck FAFEM that uses local low-pass and residual high-pass proxies to enhance deep features and produce a compact descriptor. The descriptor can then provide image-dependent information for guiding later multi-scale processing rather than being used only at the bottleneck.

This observation motivates the central idea of this work. Instead of treating bottleneck frequency enhancement and decoder multi-scale refinement as independent operations, we establish an explicit cross-level interaction between them. The frequency descriptor produced by a bottleneck FAFEM is propagated to the first decoder stage and used to regulate the relative contributions of the multi-scale branches in an MSCB. In this way, the receptive-field mixture of the decoder becomes sample dependent rather than fixed across all images.

Directly replacing the original MSCB aggregation with a fully frequency-guided formulation could, however, unnecessarily perturb a stable multi-scale representation. To address this issue, we introduce the bounded residual guidance strategy formalized in Eq. (23). Here, M_{ord} denotes the ordinary multi-scale aggregation, M_{guide} denotes the frequency-conditioned aggregation, and s is a bounded learnable guidance-strength parameter. This formulation preserves the conventional MSCB pathway as a reference and allows the model to progressively learn frequency-dependent deviations during optimization.

Based on this mechanism, we develop a ConvNeXt-Tiny U-shaped segmentation framework in which the frequency-guided modification is localized to the first encoder–decoder fusion stage; the remaining decoder stages retain conventional fusion and depthwise-separable convolution.

The principal contributions of this work are summarized as follows:

- We introduce RFG-MSCB, which uses the bottleneck FAFEM descriptor to modulate the multi-kernel MSCB branches through bounded residual interpolation around ordinary MSCB aggregation.
- We evaluate the resulting ConvNeXt-Tiny segmentation architecture on five public colorectal polyp datasets using a three-seed comparison of Baseline, FAFEM + MSCB, and Residual RFG-MSCB configurations.
- We quantify the architectural trade-off: the localized RFG-MSCB modification adds 0.08 M parameters over FAFEM + MSCB without additional reported GFLOPs at the displayed precision.

Related work

Encoder–decoder architectures for polyp segmentation

Fully convolutional encoder–decoder networks remain a central design pattern for biomedical segmentation. U-Net combines semantic context with spatial detail through skip connections, while UNet++ uses nested dense pathways to reduce the semantic gap between encoder and decoder representations [1, 2]. Polyp-specific convolutional networks subsequently developed this pattern in different directions: HardNet-MSEG couples a low-memory-traffic backbone with lightweight decoding, whereas PraNet combines a parallel partial decoder with reverse attention for progressive region refinement [7, 8]. More recent designs emphasize cross-attention and feature exploration, multi-level lesion awareness, progressive coarse-to-fine decoding, or dual-view feature propagation [9, 10, 11, 12]. These studies demonstrate the value of cross-level reuse and staged refinement for heterogeneous polyp appearance. In contrast, our formulation uses a deep frequency descriptor to control the receptive-field mixture at an earlier decoder fusion.

Modern convolutional encoders offer a complementary way to strengthen this design without replacing the full hierarchy by self-attention. ConvNeXt modernizes a ConvNet using large-kernel depthwise convolution, inverted channel expansion, Layer Normalization, GELU, layer scaling, and stochastic depth [5]. Its hierarchical outputs are suitable for U-shaped feature reuse, while depthwise spatial operators also support economical dense prediction. We use ConvNeXt-Tiny as an established encoder and focus on coupling bottleneck frequency processing with Stage-3 decoder refinement.

Transformer and hybrid contextual segmentation

Transformer and hybrid models address the broad variation in polyp size, morphology, and contrast by combining local appearance with long-range context. TransUNet provides an influential CNN–transformer formulation for medical segmentation, and Polyp-PVT adapts hierarchical pyramid-transformer features to polyp segmentation [3, 4]. Recent polyp methods extend this direction through contrastive transformer representations in CTNet, lightweight vision transformers in Polyp-LVT, contextual information flow in CIFFormer, and phase-wise feature-pyramid retention in PFPRNet [13, 14, 15, 16]. Multi-feature pyramid fusion has likewise been combined with visual transformers in MF-Net [17]. Collectively, these methods prioritize global dependency modeling and cross-scale semantic exchange, but often employ attention or transformer machinery throughout substantial parts of the network. Our method instead retains a convolutional ConvNeXt encoder–decoder and introduces one lightweight, descriptor- conditioned multi-scale refinement at Stage 3.

Efficient multi-scale convolutional decoding

Multi-scale decoding is particularly relevant to polyps, which may be small and low-contrast or occupy much of the endoscopic frame. Parallel kernels provide different effective receptive fields, while depthwise convolution limits the cost of spatial filtering. EMCAD formalized this principle in an efficient medical-image decoder whose Multi-Scale Convolution Blocks (MSCBs) expand the feature, process it with multiple depthwise kernels, aggregate the branch responses, shuffle channels, and project the result [6]. We adopt the lightweight MSCB topology as the ordinary multi-scale reference path rather than presenting it as a new block.

The broader adaptive-convolution literature shows that receptive fields need not be combined uniformly. Selective Kernel Networks use a feature descriptor to select among kernel scales, and Dynamic Convolution forms input-dependent mixtures over convolution kernels [18, 19]. Unlike these same-level routing mechanisms, RFG-MSCB transfers a low/high-frequency descriptor from the bottleneck to the Stage-3 decoder fusion. This cross-level signal weights the adopted MSCB’s 1×1 , 3×3 , and 5×5 depthwise branches while retaining their ordinary sum as the reference response. Recent medical segmentation designs that combine frequency enhancement with multi-scale adaptive attention further support the utility of jointly modeling spectral content and scale selection [20].

Frequency-aware and adaptive feature recalibration

Compact descriptors are widely used to recalibrate convolutional representations. Squeeze-and-Excitation Networks learn channel-wise gates from global pooled features, whereas FcaNet interprets channel attention in terms of selected frequency components [21, 22]. Frequency-aware processing has also moved beyond channel attention: Frequency-Aware Feature Fusion addresses the loss of high-frequency detail and inconsistency introduced when fusing features of different resolutions in dense prediction [23]. Recent medical segmentation methods have also combined frequency representations with cross-feature interaction and shape-oriented frequency partitioning [24, 25]. In polyp segmentation, Polyp-Mamba combines multi-frequency perception with gated selection, providing a domain-specific example of frequency-sensitive contextual modeling [26].

The present design differs in the source, destination, and form of its conditioning. The adopted bottleneck FAFEM produces local low-pass and residual high-pass responses and a pooled descriptor $\mathbf{z}_f$, which RFG-MSCB propagates across levels to control the Stage-3 multi-kernel mixture rather than recalibrating only the source feature. Building on frequency-aware processing and adaptive multi-scale aggregation, RFG-MSCB transfers the bottleneck descriptor across levels and uses the branch-conditioned response as a bounded residual refinement of conventional MSCB.

Boundary-aware and recent polyp-specific refinement

Weak boundaries and ambiguous foreground appearance have motivated complementary polyp-specific refinement strategies. MEGANet applies multi-scale edge-guided attention, while MSBP-Net predicts boundaries at multiple scales [27, 28]. MLB-Net combines lesion-aware and boundary-enhanced processing, and PraNet-V2 extends reverse attention with dual supervision [10, 29]. Reliability has also received

increasing attention: MEIN models multi-scale uncertainty through evidence integration rather than treating all predictions as equally certain [30]. Together with CFFormer, PFPRNet, MSPN, and DVFIIP-Net, these 2025–2026 methods illustrate a shift toward explicit contextual flow, hierarchical refinement, uncertainty modeling, and targeted boundary recovery [15, 16, 11, 12].

Boundary emphasis in our architecture is supplied by the training objective rather than a dedicated edge branch, contour decoder, or uncertainty head. The architectural focus is instead on conditioning receptive-field selection at a specific decoder level using bottleneck frequency composition while preserving an efficient multi-scale path. Accordingly, the method adopts ConvNeXt-Tiny, a locally implemented FAFEM, and the MSCB structure from EMCAD, then links the latter two through the cross-level bounded residual frequency-guided mechanism developed next.

Method

Overall Architecture

The network uses a U-shaped encoder–decoder architecture built upon a ConvNeXt-Tiny encoder. Frequency information extracted from deep semantic features regulates multi-scale processing during decoding. Frequency enhancement and multi-scale convolution are linked through a cross-level connection between the bottleneck and the first decoder stage.

Let $I \in \mathbb{R}^{3 \times H \times W}$ denote an input colonoscopy image. The ConvNeXt-Tiny encoder extracts four hierarchical feature maps,

$$\{F_1, F_2, F_3, F_4\} = E(I), \quad (1)$$

where $F_1 \in \mathbb{R}^{96 \times H/4 \times W/4}$, $F_2 \in \mathbb{R}^{192 \times H/8 \times W/8}$, $F_3 \in \mathbb{R}^{384 \times H/16 \times W/16}$, and $F_4 \in \mathbb{R}^{768 \times H/32 \times W/32}$. For an input resolution of 352×352 , their spatial dimensions are 88×88 , 44×44 , 22×22 , and 11×11 , respectively.

The deepest encoder feature F_4 is first processed by a Frequency-Aware Feature Enhancement Module (FAFEM), which produces an enhanced bottleneck representation together with a compact frequency descriptor. The enhanced representation is passed through a lightweight residual bottleneck, whereas the frequency descriptor is propagated to the first decoder stage. After the decoder feature is fused with F_3 , a Residual Frequency-Guided Multi-Scale Convolution Block (RFG-MSCB) uses the frequency descriptor to adaptively regulate its multi-scale convolution branches.

The overall processing can therefore be summarized as

$$(\tilde{F}_4, \mathbf{z}_f) = \mathcal{F}_{\text{FAFEM}}(F_4), \quad (2)$$

$$B = \mathcal{B}(\tilde{F}_4), \quad (3)$$

$$D_4 = \mathcal{M}_{\text{RFG}}(\mathcal{S}_3[\mathcal{D}_4(B), F_3], \mathbf{z}_f). \quad (4)$$

followed by

$$D_3 = \mathcal{S}_2[\mathcal{D}_3(D_4), F_2], \quad (5)$$

$$D_2 = \mathcal{S}_1[\mathcal{D}_2(D_3), F_1], \quad (6)$$

$$D_1 = \mathcal{D}_1(D_2). \quad (7)$$

Here, $\mathcal{D}_i$ denotes a decoder block, $\mathcal{S}_i$ denotes the corresponding skip-feature fusion operation, $\mathbf{z}_f$ is the frequency descriptor extracted at the bottleneck, and $\mathcal{M}_{\text{RFG}}$ denotes the residual frequency-guided multi-scale refinement mechanism. Equations (4)–(7) summarize the decoder path.

Figure 2 summarizes the complete data flow. The input image passes through the hierarchical ConvNeXt-Tiny encoder to produce F_1 – F_4 . At the bottleneck, FAFEM transforms the deepest feature F_4 into an enhanced representation and the 1536-dimensional descriptor $\mathbf{z}_f$; the enhanced representation then passes through the lightweight residual bottleneck and the fourth decoder block. After fusion with the Stage-3 encoder feature F_3 , RFG-MSCB uses $\mathbf{z}_f$ to modulate its multi-kernel branches. Thus, the descriptor is transferred from the bottleneck to a later decoder fusion rather than being used only locally. The ordinary MSCB aggregation remains the reference path, and the frequency-conditioned response enters as a bounded residual deviation. The remaining decoder stages progressively recover spatial resolution before the prediction head produces a single-channel logit map, which is bilinearly upsampled to form the final segmentation mask.




Figure 1: ConvNeXt-Tiny encoder structure. (A) Stage-level processing after the one-time ImageNet normalization: Stage 1 uses a 4×4 stride-4 stem, while subsequent stages use 2×2 stride-2 downsampling and repeat ConvNeXt blocks according to depths [3, 3, 9, 3]. (B) An individual ConvNeXt block, comprising a 7×7 stride-1 depthwise convolution, channels-last Layer Normalization, linear expansion from C to $4C$, GELU, linear projection from $4C$ to C , layer scaling, DropPath, and residual addition.

ConvNeXt-Tiny Encoder

The encoder employs ConvNeXt-Tiny with stage depths [3, 3, 9, 3] and channel dimensions [96, 192, 384, 768]. Before entering the encoder, each RGB image is normalized once using the ImageNet statistics $\mu =$




Figure 2: Overall architecture and detailed mechanism of the proposed segmentation framework. The left panel shows the ConvNeXt-Tiny encoder-decoder network with bottleneck FAFEM, Stage-3 RFG-MSCB, and the prediction head. The middle panel illustrates how FAFEM independently processes low- and high-frequency components of the deepest encoder feature, concatenates them, and generates the descriptor z_f . The right panel details RFG-MSCB, in which a lightweight guidance MLP uses z_f to adaptively weight the 1 × 1, 3 × 3, and 5 × 5 depthwise branches, and bounded residual coupling introduces the guided response as a deviation from ordinary MSCB aggregation.

(0.485, 0.456, 0.406) and $\sigma = (0.229, 0.224, 0.225)$.

Figure 1 summarizes the stage-level organization and the internal structure of an individual ConvNeXt block. The first stage begins with a 4×4 stride-4 stem convolution, whereas Stages 2–4 begin with a 2×2 stride-2 downsampling convolution. Each stage then applies its prescribed number of blocks. Within each block, a 7×7 stride-1 depthwise convolution is followed by channels-last Layer Normalization and a pointwise MLP that expands C to $4C$, applies GELU, and projects $4C$ back to C . Layer scaling and stochastic depth are applied before the transformed feature is added to the residual path.

For an input feature X , the block transformation is formalized as

$$\begin{aligned} Z &= \text{DWConv}_{7 \times 7}(X), \\ V &= W_2(\text{GELU}(W_1(\text{LN}(Z)))) . \end{aligned} \quad (8)$$

The block output is obtained through layer-scaled residual learning,

$$Y = X + \text{DropPath}(\gamma \odot V), \quad (9)$$

where γ is a learnable layer-scale parameter initialized to 10^{-6} .

A stochastic-depth rate of 0.25 is used throughout the encoder. Feature dropout with probability 0.2 is additionally applied after the second, third, and fourth encoder stages.

Frequency-Aware Bottleneck Enhancement

The deepest encoder representation F_4 contains strong semantic information but has undergone substantial spatial downsampling. To expose complementary frequency characteristics of this representation, FAFEM is applied exclusively at the bottleneck.

Given $X = F_4$, a local low-frequency approximation is obtained using 3×3 average pooling with stride 1 and padding 1,

$$X_L = \text{AvgPool}_{3 \times 3}(X), \quad (10)$$

$$X_H = X - X_L. \quad (11)$$

The low- and high-frequency components in Eqs. (10) and (11) are independently transformed using depthwise convolutional branches,

$$\hat{X}_L = \phi_L(X_L), \quad \hat{X}_H = \phi_H(X_H), \quad (12)$$

where each $\phi(\cdot)$ contains a 3×3 depthwise convolution followed by GELU activation.

The two components are concatenated and globally pooled to obtain the frequency descriptor

$$\mathbf{z}_f = \text{GAP}([\hat{X}_L, \hat{X}_H]) \in \mathbb{R}^{2C}. \quad (13)$$

Since $C = 768$ at the bottleneck, $\mathbf{z}_f \in \mathbb{R}^{1536}$.

A lightweight channel-gating network applies global average pooling to $[\hat{X}_L, \hat{X}_H]$, followed by 1×1 transformations $1536 \rightarrow 48 \rightarrow 1536$, GELU after the first transformation, and a sigmoid after the second. Splitting the output channel-wise gives $w_L, w_H \in (0, 1)^{768}$; these weights are not normalized to sum to one. Their weighted combination is spatially refined and introduced through a residual connection:

$$\tilde{F}_4 = F_4 + \gamma_f \mathcal{R}_f \left(w_L \odot \hat{X}_L + w_H \odot \hat{X}_H \right), \quad (14)$$

where $\mathcal{R}_f$ is a bias-free 3×3 depthwise convolution with stride 1 and padding 1, and γ_f is a learnable scalar initialized to zero. The descriptor $\mathbf{z}_f$ is pooled before the channel weights are applied.

Lightweight Residual Bottleneck

The frequency-enhanced representation $\tilde{F}_4$ is subsequently processed by a lightweight residual bottleneck.

A 1×1 convolution first reduces the number of channels from 768 to 192:

$$X_r = \text{GELU} \left(\text{LN} \left(\text{Conv}_{1 \times 1}(\tilde{F}_4) \right) \right). \quad (15)$$

A bias-free 3×3 depthwise-separable convolution with padding 1 performs spatial transformation in the reduced feature space, followed by channels-first Layer Normalization ($\epsilon = 10^{-6}$), GELU, and spatial dropout with probability 0.2. A bias-free 1×1 convolution then projects the representation back to 768 channels and is followed by channels-first Layer Normalization.

The bottleneck output is formulated as

$$B = \tilde{F}_4 + \gamma_b \mathcal{H}(\tilde{F}_4), \quad (16)$$

where $\mathcal{H}(\cdot)$ denotes the bottleneck transformation and γ_b is a learnable scalar initialized to zero.

This initialization makes the bottleneck behave initially as an identity mapping and allows its transformation to be progressively introduced during optimization.

Residual Frequency-Guided Multi-Scale Convolution

The principal architectural contribution of this work is the Residual Frequency-Guided Multi-Scale Convolution mechanism.

After the bottleneck feature is upsampled from $H/32$ to $H/16$, the decoder representation is reduced from 768 to 384 channels and concatenated with the Stage-3 encoder feature F_3 . A lightweight fusion block transforms the resulting 768-channel representation into $X_3 \in \mathbb{R}^{384 \times H/16 \times W/16}$, or $\mathbb{R}^{384 \times 22 \times 22}$ for a 352×352 input.

A conventional lightweight multi-scale convolution block first expands the feature dimension from C to $2C$:

$$X_e = \text{ReLU6}(\text{BN}(\text{Conv}_{1 \times 1}(X_3))). \quad (17)$$

Three parallel depthwise convolution branches with kernel sizes $k \in \{1, 3, 5\}$ then generate multi-scale responses,

$$B_k = \phi_k(X_e), \quad (18)$$

where each ϕ_k consists of a depthwise convolution, Batch Normalization, and ReLU6 activation.

In a conventional MSCB, these branches are combined uniformly:

$$M_{\text{ord}} = B_1 + B_3 + B_5. \quad (19)$$

The new mechanism conditions the relative contributions of the three receptive fields on the bottleneck frequency descriptor $\mathbf{z}_f$.

A lightweight guidance network maps the 1536-dimensional descriptor to three branch logits:

$$\mathbf{q} = W_{g2}(\text{GELU}(W_{g1}\mathbf{z}_f)), \quad (20)$$

where $W_{g1} : 1536 \rightarrow 48$ and $W_{g2} : 48 \rightarrow 3$.

The normalized branch weights are calculated as

$$[\alpha_1, \alpha_3, \alpha_5] = 3 \text{ Softmax}(\mathbf{q}), \quad (21)$$

such that $\alpha_1 + \alpha_3 + \alpha_5 = 3$.

The frequency-guided response becomes

$$M_{\text{guide}} = \alpha_1 B_1 + \alpha_3 B_3 + \alpha_5 B_5. \quad (22)$$

We introduce the guidance as a bounded residual deviation from the conventional multi-scale response:

$$M_{\text{RFG}} = M_{\text{ord}} + s(M_{\text{guide}} - M_{\text{ord}}), \quad (23)$$

where $0 \leq s \leq 0.5$.

The implementation stores an unconstrained learnable scalar θ_s , initialized to 0.05, and uses the hard-clamped effective strength $s = \text{clamp}(\theta_s, 0, 0.5)$ in Eq. (23).

The final guidance-layer weights and biases are initialized so that the branch logits are equal. Thus, at initialization, $\alpha_1 = \alpha_3 = \alpha_5 = 1$ and $M_{\text{guide}} = M_{\text{ord}}$.

The module thus initially reproduces the conventional multi-scale path, while optimization is allowed to learn image-dependent deviations from uniform branch weighting.

The resulting feature is processed using channel shuffle with $\text{gcd}(2C, C) = 384$ groups and a bias-free 1×1 projection from $2C$ back to C , followed by Batch Normalization. This sequence defines $\mathcal{P}$ below. A final residual connection preserves the input representation:

$$Y_3 = X_3 + \mathcal{P}(\text{Shuffle}(M_{\text{RFG}})). \quad (24)$$

This formulation connects the bottleneck frequency descriptor to multi-scale spatial refinement at $H/16$ resolution.

Progressive Decoder and Skip Fusion

After frequency-guided Stage-3 refinement, the decoder progressively restores spatial resolution using bilinear upsampling and depthwise-separable convolution.

The first refined decoder feature has 384 channels, $D_4 = Y_3$.

The subsequent decoder stages follow Eqs. (5)–(7).

For a 352×352 input, the decoder dimensions are $D_4 : 384 \times 22 \times 22$, $D_3 : 192 \times 44 \times 44$, $D_2 : 96 \times 88 \times 88$, and $D_1 : 96 \times 176 \times 176$.

At each skip level, the decoder feature is resized to the skip-feature shape when needed using bilinear interpolation with `align_corners=False`, concatenated with the corresponding encoder feature, and processed by a lightweight fusion block. The fusion block contains a bias-free 1×1 projection, channels-first Layer Normalization ($\epsilon = 10^{-6}$), GELU, a bias-free depthwise-separable 3×3 convolution with padding 1, a second

channels-first Layer Normalization and GELU, and spatial dropout with probability 0.2. The primary twofold upsampling inside each decoder block uses bilinear interpolation with `align_corners=True`.

No additional attention mechanism, frequency module, MSCB, or boundary-specific refinement block is applied at these decoder stages.

The final decoder stage has no direct encoder skip connection because the ConvNeXt stem begins at one-quarter of the input resolution.

Prediction Head

The final decoder feature D_1 is processed using a lightweight refinement head. Two depthwise-separable 3×3 convolutions transform the feature dimensions as $96 \rightarrow 96 \rightarrow 48$.

Each transformation is followed by Layer Normalization and GELU activation. A final 1×1 convolution produces a single-channel segmentation logit map:

$$P_0 = \text{Conv}_{1 \times 1}(\mathcal{R}(D_1)), \quad (25)$$

$$P = \text{Up}_{H,W}(P_0), \quad (26)$$

$$\hat{Y} = \sigma(P). \quad (27)$$

The prediction in Eq. (26) is bilinearly resized from $H/2 \times W/2$ to the original input resolution with `align_corners=False`. The network directly returns logits; during inference, Eq. (27) gives the foreground probability map,

where $\sigma(\cdot)$ denotes the sigmoid function.

The selected architecture does not employ auxiliary prediction heads or deep supervision.

Model Complexity

The complete model has 30.06 M parameters and requires 27.82 GFLOPs for a 352×352 input. The FAFEM + MSCB configuration has 29.98 M parameters and 27.818 GFLOPs; the final RFG modification therefore adds approximately 0.08 M parameters, while the GFLOP estimate remains 27.818 at the displayed precision. Adaptive computation is localized to the Stage-3 encoder-decoder fusion, and the remaining decoder stages retain depthwise-separable convolution. Latency, throughput, and peak memory were not retained with the final runs, so these measurements support a complexity comparison but not a matched-hardware runtime claim.

Table 1: Computational complexity comparison. Parameter counts and computational cost are taken from original publications or official implementations where available; ‘–’ denotes an unavailable or unverified value.

Method	Year	Input	Params (M)	GFLOPs
MEGANet (Res2Net-50) [27]	2024	352^2	44.19	28.85
MSBP-Net [28]	2026	352^2	25.52	12.86
CIFFormer [15]	2025	352^2	32.44	13.14
PFPRNet [16]	2025	352^2	118.52	39.39
MLB-Net [10]	2026	352^2	87.70	–
MSPN [11]	2026	352^2	28.56	13.21
DVFIP-Net [12]	2026	352^2	45.62	19.37
Ours	–	352^2	30.06	27.82 [†]

[†]Our value uses 13.91 GMACs and the explicit convention $\text{FLOPs} = 2 \times \text{MACs}$. Competitor values are preserved in source terminology and are not convention-normalized.

Experiments

Datasets and Data Splits

To evaluate the proposed framework under a fixed training and evaluation design, we used Kvasir-SEG [31], CVC-ClinicDB [32], CVC-300 [33], CVC-ColonDB [34], and ETIS-LaribPolypDB [35]. Following a PraNet-style data partition, 100 of the 1000 Kvasir-SEG images and 62 of the 612 CVC-ClinicDB images were reserved as final seen test sets. The remaining 900 Kvasir-SEG and 550 CVC-ClinicDB images formed a combined 1450-image train/validation pool. This pool was split deterministically into 1305 training images and 145 internal validation images using shuffled sampling with `random_state=42`; dataset membership was identical across all training seeds.

The final Kvasir-SEG and CVC-ClinicDB test images were excluded from training, validation, checkpoint selection, early stopping, threshold tuning, hyperparameter tuning, and architecture selection. CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB were used exclusively as external test sets and did not influence model selection.

Table 2: Dataset partition used in the experiments.

Dataset	Total	Train	Test
Kvasir-SEG	1000	900	100
CVC-ClinicDB	612	550	62
CVC-300	60	–	60
CVC-ColonDB	380	–	380
ETIS-LaribPolypDB	196	–	196

Images were converted to RGB, resized to 352×352 with bilinear interpolation, and scaled to $[0, 1]$. Masks were resized with nearest-neighbor interpolation and binarized. The encoder applied ImageNet normalization with mean $(0.485, 0.456, 0.406)$ and standard deviation $(0.229, 0.224, 0.225)$.

Implementation Details

The final protocol uses seeds 42, 6543, and 7777. All final and ablation results are reported as the mean $\pm$ sample standard deviation (ddof=1) across these seeds unless otherwise noted. The runs were implemented in PyTorch 2.5.1+cu124 on Windows and executed on an NVIDIA GeForce RTX 4090, with a maximum of 200 epochs and batch size 24. AdamW with betas $(0.9, 0.999)$ used separate parameter groups. The maximum learning rates were 3.0×10^{-4} for the decoder and 4.5×10^{-4} for the final refinement head. Encoder Stages 4, 3, 2, and 1 plus the stem used 3.0×10^{-5} , 2.4×10^{-5} , 1.92×10^{-5} , and 1.536×10^{-5} , respectively, corresponding to layer-wise decay 0.8. Weight decay was 1.0×10^{-4} for the decoder, 5.0×10^{-2} for the encoder, and 1.0×10^{-2} for the refinement head.

Learning rates were warmed up linearly for five epochs from 0.1 to 1.0 of each group maximum and then cosine-annealed over the remaining epochs to 1.0×10^{-6} . Training used AMP FP16, global gradient clipping at norm 1.0, and an exponential moving average (EMA) with decay 0.995. All parameters were trainable from the first epoch; no staged freezing or decoder-only warmup was used. The encoder was initialized from `convnext_tiny_22k_1k_384.pth`, logged as ImageNet pretrained weights. Deep supervision and frequency augmentation were disabled, and sampling was uniform. Early stopping was configured with patience 30 from the start of training. Seeds 42 and 6543 completed the maximum 200 epochs, corresponding to zero-based epochs 0–199; seed 7777 stopped after epoch 167 (168 completed epochs). Model complexity and its measurement convention are reported in Table 1.

Training augmentation applied horizontal and vertical flips independently with probability 0.5, rotation in $[-15^\circ, 15^\circ]$ with probability 0.5, and a random crop ratio in $[0.7, 1.0]$ followed by resizing with probability 0.5. Joint brightness and contrast factors in $[0.85, 1.15]$ and saturation in $[0.9, 1.1]$ were applied with probability 0.5. A 3×3 Gaussian blur and additive Gaussian noise with standard deviation 0.015 were applied with probabilities 0.15 and 0.25, respectively.

The code seeded Python, NumPy, PyTorch CPU, and all CUDA random-number generators. Deterministic CUDA algorithms and cuDNN deterministic backend flags were not enabled.

Loss Function

The final objective was

$$\mathcal{L} = 0.55\mathcal{L}_{\text{Dice}} + 0.25\mathcal{L}_{\text{BCE}} + 0.20\mathcal{L}_{\text{boundary}}. \quad (28)$$

The objective in Eq. (28) uses a soft Dice term with a smoothing constant of 1.0. The BCE-with-logits term used targets smoothed as $y_{\text{BCE}} = 0.98y + 0.01$, corresponding to label smoothing 0.02. For the boundary-aware term, fixed Sobel filters were applied to each binary mask to obtain $g = \sqrt{g_x^2 + g_y^2}$. The magnitude was normalized per image as $a = g / \max(\max(g), 10^{-6})$, and the pixel weights were $w = 1 + \kappa a$ with $\kappa = 5$; consequently, the weights range from 1 to 6. For probability $p = \sigma(P)$ and target y , the term was

$$\mathcal{L}_{\text{boundary}} = -\frac{\sum_i w_i [y_i \log(p_i + \epsilon) + (1 - y_i) \log(1 - p_i + \epsilon)]}{\sum_i w_i + \epsilon} + 1 - \frac{\sum_i w_i y_i p_i + \epsilon}{\sum_i w_i (y_i + p_i - y_i p_i) + \epsilon}, \quad (29)$$

with $\epsilon = 10^{-6}$ and batch averaging for the weighted IoU component. The focal-Tversky term was assigned zero weight and therefore did not contribute to the final objective.

Evaluation Metrics

We report mean Dice (mDice), mean intersection over union (mIoU), S-measure S_α [36], weighted F-measure F_β^w [37], mean and maximum E-measure (mE_ϕ and $\max E_\phi$) [38], and mean absolute error (MAE) [39]. For each image, the thresholded prediction P and binary ground truth G were flattened and evaluated as

$$\text{Dice}(P, G) = \frac{2|P \cap G| + 10^{-6}}{|P| + |G| + 10^{-6}}, \quad \text{IoU}(P, G) = \frac{|P \cap G| + 10^{-6}}{|P| + |G| - |P \cap G| + 10^{-6}}. \quad (30)$$

The reported mDice and mIoU are arithmetic means of the per-image scores in Eq. (30), rather than global pixel aggregates. S-measure, weighted F-measure, the E-measure curve, and MAE were computed per image with `py_sod_metrics` from 8-bit probability maps and binary ground truth. The mean and maximum of the resulting E-measure curve were reported.

Model Selection and Evaluation Protocol

At each epoch, EMA weights were evaluated only on the 145-image internal validation subset. Probabilities were obtained with one sigmoid operation and binarized using the strict rule $p > 0.45$. A checkpoint was saved only when internal-validation IoU was strictly greater than the previous best. Best checkpoints were selected independently for each seed according to internal-validation IoU.

After model selection, each best checkpoint was loaded strictly and evaluated on the 100-image Kvasir-SEG and 62-image CVC-ClinicDB seen test sets, followed by the full CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB external test sets. Test-time augmentation was disabled. Neither the seen nor external test results influenced checkpoint selection or other model-development decisions.

Experimental Study

The experimental study first examines the contribution of the architectural components under controlled variants and then compares the resulting model with published methods. Presenting the ablation before the comparative evaluation establishes the reference configuration and clarifies which part of the final design is being assessed in the comparative tables.

Ablation Study

The ablation design compares three configurations in sequence: Baseline, FAFEM + MSCB, and Residual RFG-MSCB (Ours). This sequence evaluates the feature-enhancement and multi-scale components together before assessing the additional residual frequency-guided modification.

Ablation Results

The Baseline serves as the reference configuration without FAFEM and MSCB. Adding FAFEM and MSCB improves aggregate segmentation performance across the five evaluation datasets. Residual RFG-MSCB then introduces frequency-guided coupling, allowing the bottleneck descriptor to modulate the multi-scale branches at the Stage-3 decoder fusion. Table 3 summarizes the arithmetic mean of each segmentation metric across Kvasir-SEG, CVC-ClinicDB, CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB together with the corresponding model complexity.

Table 3: Ablation comparison averaged across the five polyp segmentation datasets, with computational complexity.

Model	mDice	mIoU	F_{β}^w	S_{α}	mE_{ϕ}	$\max E_{\phi}$	MAE	Params (M)	GFLOPs
Baseline	0.827	0.763	0.811	0.890	0.907	0.913	0.019	29.19	27.202
FAFEM + MSCB	<u>0.869</u>	<u>0.803</u>	<u>0.845</u>	<u>0.910</u>	<u>0.937</u>	<u>0.947</u>	<u>0.018</u>	29.98	27.818
Residual RFG-MSCB (Ours)	0.887	0.832	0.864	0.923	0.952	0.956	0.014	30.06	27.818

Ablation Interpretation

Across the five-dataset averages, FAFEM + MSCB increases mDice, mIoU, F_{β}^w , S_{α} , mE_{ϕ} , and $\max E_{\phi}$ over the Baseline by 0.042, 0.040, 0.034, 0.020, 0.030, and 0.034, respectively, corresponding to gains of 4.2, 4.0, 3.4, 2.0, 3.0, and 3.4 percentage points. It reduces MAE by 0.001. Residual RFG-MSCB further increases these six metrics over FAFEM + MSCB by 0.018, 0.029, 0.019, 0.013, 0.015, and 0.009, respectively (1.8, 2.9, 1.9, 1.3, 1.5, and 0.9 percentage points), and reduces MAE by 0.004. The parameter count increases from 29.19 M to 29.98 M and then 30.06 M, changes of 0.79 M and 0.08 M, while GFLOPs increase from 27.202 to 27.818 and then remain unchanged, changes of 0.616 and 0.000. These are descriptive comparisons; no statistical significance is claimed.

The qualitative examples in Fig. 3 complement the aggregate ablation results in Table 3 by illustrating how the architectural changes affect predicted mask geometry in the displayed samples.

Qualitative Ablation Analysis

Relative to the Baseline, FAFEM + MSCB produces mask shapes with closer spatial coverage or boundary alignment in several displayed examples. In the displayed CVC-ColonDB case with an elongated annotation, the disconnected Baseline regions are replaced by a more coherent prediction. For the displayed small-lesion ETIS-LaribPolypDB example, the Baseline does not recover the annotated region, whereas both augmented configurations produce a localized response.

Residual RFG-MSCB yields a more coherent or better-localized mask in selected cases, including the displayed CVC-ColonDB case, while its effect varies among the selected images. In the displayed Kvasir-SEG, CVC-ClinicDB, and CVC-300 examples, the predictions differ mainly in boundary position and extent, and no configuration reproduces every ground-truth contour exactly. These examples therefore provide a visual complement to the aggregate trends in Table 3, rather than evidence of uniform improvement for every image.

Stage Placement

The selected configuration localizes RFG-MSCB to Stage 3. Systematic multi-seed evaluation of alternative placements remains future work.




Figure 3: Qualitative comparison of the three ablation configurations using two displayed samples per dataset. Dataset groups correspond to Kvasir-SEG, CVC-ClinicDB, CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB; columns show the original image, ground-truth mask, Baseline prediction, FAFEM + MSCB prediction, and Residual RFG-MSCB prediction.

Comparative study

After the component-level analysis, the comparative study evaluates the complete model on the Kvasir-SEG and CVC-ClinicDB seen test sets and on three external test sets, and places its performance alongside published baselines. The quantitative results are presented below.

Quantitative analysis

On Kvasir-SEG, the proposed model achieved the highest displayed mDice and mIoU, reaching 0.932 ± 0.0044 and 0.883 ± 0.004 , respectively. PFPRNet was second-highest for both metrics, with mDice of 0.930 and mIoU of 0.881. On CVC-ClinicDB, the proposed model achieved 0.948 ± 0.003 mDice, which was second-highest to PFPRNet’s 0.949, and the highest displayed mIoU at 0.904 ± 0.004 . It also tied for the highest displayed S_α on CVC-ClinicDB at 0.958 ± 0.001 and was second-highest in mE_ϕ and $\max E_\phi$ at 0.979 ± 0.001 and 0.988 ± 0.002 , respectively.

The main comparison is restricted to methods using comparable PraNet-style data protocols with closely matched Kvasir-SEG and CVC-ClinicDB partitioning. Methods using materially different data partitions remain discussed in Related Work but are omitted from the main performance tables. Differences in optimization, augmentation, backbone initialization, model-selection criteria, and implementation details remain. Across the comparison tables, the best and second-best displayed values are shown in bold and underlined, respectively.

On the external datasets, mDice/mIoU were $0.916 \pm 0.002/0.863 \pm 0.002$ on CVC-300, $0.824 \pm 0.009/0.747 \pm 0.008$ on CVC-ColonDB, and $0.816 \pm 0.012/0.763 \pm 0.013$ on ETIS-LaribPolypDB. Across these three external datasets, the model’s mDice ranged from 0.816 to 0.916 and its mIoU ranged from 0.747 to 0.863; the largest between-seed variation in both metrics occurred on ETIS. On CVC-300, the model achieved the highest displayed mDice, mIoU, S_α , and $\max E_\phi$ and tied for the lowest displayed MAE. Its mIoU, mE_ϕ , and $\max E_\phi$ were second-highest on ETIS, while its mIoU and $\max E_\phi$ were second-highest and its MAE was second-lowest on CVC-ColonDB.

Table 4: Quantitative comparison on Kvasir-SEG.

Methods	Year	mDice	mIoU	F_β^w	S_α	mE_ϕ	$\max E_\phi$	MAE
MEGANet [27]	2024	0.913	0.863	0.907	0.918	0.956	0.959	0.025
MSBP-Net [28]	2026	0.919	0.868	0.916	0.926	0.961	–	0.023
CIFFormer [15]	2025	0.926	0.876	0.929	0.939	0.967	–	0.019
PFPRNet [16]	2025	<u>0.930</u>	<u>0.881</u>	0.920	0.935	0.971	–	0.019
MLB-Net [10]	2026	0.926	0.878	<u>0.921</u>	<u>0.938</u>	<u>0.969</u>	0.975	0.021
MSPN [11]	2026	0.927	0.879	–	–	–	–	0.021
DVFIP-Net [12]	2026	0.922	0.874	–	–	–	–	–
Ours	–	0.932	0.883	0.914	0.936	0.962	0.974	0.020
		± 0.0044	± 0.004	± 0.006	± 0.003	± 0.003	± 0.002	± 0.001

Table 5: Quantitative comparison on CVC-ClinicDB.

Methods	Year	mDice	mIoU	F_β^w	S_α	mE_ϕ	$\max E_\phi$	MAE
MEGANet [27]	2024	0.938	0.894	0.940	0.950	–	0.987	0.007
MSBP-Net [28]	2026	0.940	0.892	0.927	0.939	0.976	–	0.007
CIFFormer [15]	2025	0.944	0.897	0.941	0.934	0.972	–	0.007
PFPRNet [16]	2025	0.949	<u>0.903</u>	<u>0.947</u>	<u>0.957</u>	0.989	–	0.006
MLB-Net [10]	2026	0.945	0.902	0.948	0.958	0.989	0.992	0.006
MSPN [11]	2026	0.943	0.902	–	–	–	–	0.006

Continued on next page

Table 5: Quantitative comparison on CVC-ClinicDB (continued).

Methods	Year	mDice	mIoU	F_{β}^w	S_{α}	mE_{ϕ}	$\max E_{\phi}$	MAE
DVFIP-Net [12]	2026	0.942	0.896	—	—	—	—	—
Ours	—	0.948 ± 0.003	0.904 ± 0.004	0.939 ± 0.001	0.958 ± 0.001	0.979 ± 0.001	0.988 ± 0.002	0.007 ± 0.001

Table 6: Quantitative comparison on ETIS-LaribPolypDB.

Methods	Year	mDice	mIoU	F_{β}^w	S_{α}	mE_{ϕ}	$\max E_{\phi}$	MAE
MEGANet [27]	2024	0.789	0.709	0.753	0.866	0.855	0.915	0.015
MSBP-Net [28]	2026	0.795	0.718	0.770	0.873	0.899	—	0.019
CIFFormer [15]	2025	0.810	0.746	0.858	0.889	0.918	—	<u>0.013</u>
PFPRNet [16]	2025	0.828	0.764	<u>0.839</u>	0.908	0.935	—	0.012
MLB-Net [10]	2026	<u>0.821</u>	0.744	0.780	<u>0.899</u>	0.914	0.935	<u>0.013</u>
MSPN [11]	2026	0.798	0.717	—	—	—	—	0.017
DVFIP-Net [12]	2026	0.811	0.730	—	—	—	—	—
Ours	—	0.816 ± 0.012	0.763 ± 0.013	0.783 ± 0.004	0.882 ± 0.001	0.921 ± 0.006	0.927 ± 0.010	0.014 ± 0.001

Table 7: Quantitative comparison on CVC-ColonDB.

Methods	Year	mDice	mIoU	F_{β}^w	S_{α}	mE_{ϕ}	$\max E_{\phi}$	MAE
MEGANet [27]	2024	0.793	0.714	0.779	0.854	0.892	0.895	0.034
MSBP-Net [28]	2026	0.810	0.731	0.810	0.861	0.912	—	0.029
CIFFormer [15]	2025	0.823	0.741	0.851	0.874	<u>0.918</u>	—	0.026
PFPRNet [16]	2025	0.819	0.741	<u>0.843</u>	<u>0.882</u>	0.915	—	0.028
MLB-Net [10]	2026	0.832	0.757	0.825	0.886	0.924	0.931	0.023
MSPN [11]	2026	0.811	0.734	—	—	—	—	0.029
DVFIP-Net [12]	2026	<u>0.826</u>	0.746	—	—	—	—	—
Ours	—	0.824 ± 0.009	0.747 ± 0.008	0.832 ± 0.002	0.862 ± 0.004	0.911 ± 0.003	0.919 ± 0.002	0.025 ± 0.001

Table 8: Quantitative comparison on CVC-300 (EndoScene).

Methods	Year	mDice	mIoU	F_{β}^w	S_{α}	mE_{ϕ}	$\max E_{\phi}$	MAE
MEGANet [27]	2024	0.899	0.834	0.882	0.935	0.966	0.969	0.007
MSBP-Net [28]	2026	0.903	<u>0.859</u>	0.927	0.942	0.967	—	0.019
CIFFormer [15]	2025	0.883	0.822	0.912	0.939	0.961	—	0.100
PFPRNet [16]	2025	0.891	0.833	<u>0.917</u>	0.928	0.950	—	0.013
MLB-Net [10]	2026	<u>0.915</u>	0.852	0.909	<u>0.974</u>	0.974	<u>0.978</u>	0.005
MSPN [11]	2026	0.907	0.840	—	—	—	—	<u>0.006</u>
DVFIP-Net [12]	2026	0.900	0.833	—	—	—	—	—
Ours	—	0.916 ± 0.002	0.863 ± 0.002	0.847 ± 0.005	0.975 ± 0.004	0.969 ± 0.004	0.979 ± 0.005	0.005 ± 0.001

Discussion

Quantitative and Ablation Interpretation

The evaluation protocol separates seen-domain testing on Kvasir-SEG and CVC-ClinicDB from external testing on CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB. All five test sets are isolated from checkpoint selection, which uses only the internal validation subset of the combined development pool.

The controlled ablation evaluates Baseline, FAFEM + MSCB, and Residual RFG-MSCB using arithmetic means across the five datasets. The final increment over FAFEM + MSCB increases mDice from 0.869 to 0.887 and mIoU from 0.803 to 0.832 while reducing MAE from 0.018 to 0.014. These aggregate changes describe the combined effect of the residual frequency-guided formulation at the Stage-3 fusion.

The bounded residual design preserves the conventional MSCB response while allowing sample-dependent branch modulation. Because residualization and frequency conditioning were not evaluated independently, their individual contributions remain unresolved. This adaptive refinement reuses the bottleneck descriptor at one decoder stage without discarding the established multi-scale response or modifying the full decoder.

Cross-Dataset Generalization

CVC-300, CVC-ColonDB, and ETIS-LaribPolypDB are reserved for external generalization evaluation and do not influence model development. Differences in acquisition, annotation, case mix, and polyp morphology make these datasets a distinct test of transfer beyond the Kvasir-SEG/CVC-ClinicDB development domain.

On the external datasets, the model obtained mDice/mIoU of 0.916/0.863 on CVC-300, 0.824/0.747 on CVC-ColonDB, and 0.816/0.763 on ETIS-LaribPolypDB. Performance varied across these datasets, with the strongest transfer observed on CVC-300. On CVC-ColonDB and ETIS, the model’s mIoU was second-highest among the displayed methods, whereas its mDice was below the two highest displayed values. Differences in acquisition, annotation, and lesion appearance are plausible contributors to the lower CVC-ColonDB and ETIS results, but were not isolated experimentally.

Although the main comparison is restricted to methods with closely matched dataset partitioning, differences in optimization, augmentation, backbone initialization, model-selection criteria, and implementation details remain.

Practical Trade-offs and Placement

As summarized in Table 1, the final architecture has 30.06 M parameters and requires 27.82 GFLOPs at 352×352 . Matched-hardware latency was not measured. The final model localizes the new guidance to the Stage-3 decoder–encoder fusion. Systematic multi-seed evaluation of alternative placements and deployment-oriented latency optimization remain future work.

Limitations

- **No statistical significance testing:** All reported differences are descriptive; no hypothesis testing was performed.
- **Seed coverage:** The controlled comparisons use three training seeds, which limits characterization of run-to-run variability.
- **Component isolation:** Residualization and frequency conditioning were not independently disentangled.
- **Placement evaluation:** Alternative RFG-MSCB placements were not compared in a controlled multi-seed study.
- **Runtime evaluation:** No matched-hardware latency or peak-memory benchmark was retained.
- **External generalization:** Performance was lower on CVC-ColonDB and ETIS-LaribPolypDB than on CVC-300, motivating broader evaluation across acquisition domains.

Future Work

- Systematic multi-seed placement ablation with controlled batch size and learning rate.
- Disentanglement of residualization, frequency conditioning, and initialization through factorial experiments.
- Domain adaptation or targeted augmentation for CVC-ColonDB and ETIS-LaribPolypDB.
- Evaluation of between-configuration differences using larger seed ensembles or bootstrap confidence intervals.

Conclusion

We presented a ConvNeXt-Tiny U-shaped segmentation framework that couples a locally implemented bottleneck FAFEM descriptor to MSCB through RFG-MSCB at the Stage-3 encoder–decoder fusion. The design extends frequency enhancement into decoder refinement while preserving the ordinary multi-scale pathway as a reference.

The final model achieved mDice/mIoU of 0.932/0.883 on Kvasir-SEG and transferred most strongly to CVC-300 among the three external datasets, with lower performance on CVC-ColonDB and ETIS-LaribPolypDB. Across the five-dataset aggregate, Residual RFG-MSCB increased mDice from 0.869 to 0.887 over FAFEM + MSCB. The localized modification adds 0.08 M parameters without additional reported GFLOPs at the displayed precision and leaves the remaining depthwise-separable decoder stages unchanged.

Future work should evaluate alternative placements across multiple seeds, disentangle residualization from frequency conditioning, and examine adaptation across acquisition domains.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The datasets used in this study are publicly available from the sources cited in the manuscript. The experiment implementation is available at <https://github.com/shosseinj/convnext-unet>.

References

- [1] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, pages 234–241. Springer, 2015.
- [2] Zongwei Zhou, Md Mahfuzur Rahman Siddiquee, Nima Tajbakhsh, and Jianming Liang. UNet++: Redesigning skip connections to exploit multiscale features in image segmentation. *IEEE Transactions on Medical Imaging*, 39(6):1856–1867, 2020.
- [3] Jieneng Chen, Yongyi Lu, Qihang Yu, Xiangde Luo, Ehsan Adeli, Yan Wang, Le Lu, Alan L. Yuille, and Yuyin Zhou. TransUNet: Transformers make strong encoders for medical image segmentation, 2021. arXiv preprint arXiv:2102.04306.
- [4] Bo Dong, Wenhai Wang, Deng-Ping Fan, Jinpeng Li, Huazhu Fu, and Ling Shao. Polyp-PVT: Polyp segmentation with pyramid vision transformers. *CAAI Artificial Intelligence Research*, page 9150015, 2023.
- [5] Zhuang Liu, Hanzi Mao, Chao-Yuan Wu, Christoph Feichtenhofer, Trevor Darrell, and Saining Xie. A convnet for the 2020s. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11966–11976, 2022.
- [6] Md Mostafijur Rahman, Mustafa Munir, and Radu Marculescu. EMCAD: Efficient multi-scale convolutional attention decoding for medical image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11769–11779, 2024.
- [7] Chien-Hsiang Huang, Hung-Yu Wu, and Youn-Long Lin. HardNet-MSEG: A simple encoder-decoder polyp segmentation neural network that achieves over 0.9 mean dice and 86 fps, 2021. arXiv preprint arXiv:2101.07172.
- [8] Deng-Ping Fan, Ge-Peng Ji, Tao Zhou, Geng Chen, Huazhu Fu, Jianbing Shen, and Ling Shao. PraNet: Parallel reverse attention network for polyp segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2020*, pages 263–273. Springer, 2020.
- [9] Guoqi Liu, Sheng Yao, Dong Liu, Baofang Chang, Zongyu Chen, Jiajia Wang, and Jiangqi Wei. CAFE-Net: Cross-attention and feature exploration network for polyp segmentation. *Expert Systems with Applications*, 238:121754, 2024.
- [10] Juntong Ti, Lijun Liu, Xiaobing Yang, Li Liu, and Wei Peng. MLB-Net: A multi-level lesion-aware and boundary-enhanced network for polyp segmentation. *Journal of Imaging Informatics in Medicine*, 2026.

- [11] Yi Yang, Junlong Cheng, Taotao Zhu, and Min Zhu. From coarse to fine: Multi-stage progressive network for colon polyp segmentation. In *2026 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 7652–7656, 2026.
- [12] Haijiao Yun, Chao Wang, Jinyang Luan, Ziqing Han, Qingyu Du, and Mingjing Li. DVFIP-Net: Dual view feature interaction propagation network for polyp segmentation. *IEEE Access*, 14:31995–32008, 2026.
- [13] Bin Xiao, Jinwu Hu, Weisheng Li, Chi-Man Pun, and Xiuli Bi. CTNet: Contrastive transformer network for polyp segmentation. *IEEE Transactions on Cybernetics*, 54:5040–5053, 2024.
- [14] Long Lin, Guangzu Lv, Bin Wang, Cunlu Xu, and Jun Liu. Polyp-LVT: Polyp segmentation with lightweight vision transformers. *Knowledge-Based Systems*, 300:112181, 2024.
- [15] Cunlu Xu, Long Lin, Bin Wang, and Jun Liu. CFFormer: A contextual information flow guided transformer for colorectal polyp segmentation. *Neurocomputing*, 644:130413, 2025.
- [16] Jinghui Chu, Wangtao Liu, Qi Tian, and Wei Lu. PFPRNet: A phase-wise feature pyramid with retention network for polyp segmentation. *IEEE Journal of Biomedical and Health Informatics*, 29:1137–1150, 2025.
- [17] Shuangyuan Wang, Shengmao Lin, Fujia Sun, and Xiaobo Li. Multi-feature fusion for accurate polyp segmentation using pyramid visual transformers. *Expert Systems with Applications*, 280:127558, 2025.
- [18] Xiang Li, Wenhai Wang, Xiaolin Hu, and Jian Yang. Selective kernel networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 510–519, 2019.
- [19] Yinpeng Chen, Xiyang Dai, Mengchen Liu, Dongdong Chen, Lu Yuan, and Zicheng Liu. Dynamic convolution: Attention over convolution kernels. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11027–11036, 2020.
- [20] Ke Peng, Bi Hong, Yiyang Liu, Chunyu Chen, Yu Pan, Yu Jiang, and Xiaojuan Liu. FE-MAANet: A frequency-enhanced multi-scale adaptive attention network for medical image segmentation. *Biomedical Signal Processing and Control*, 113:108884, 2026.
- [21] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7132–7141, 2018.
- [22] Zequn Qin, Pengyi Zhang, Fei Wu, and Xi Li. FcaNet: Frequency channel attention networks. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 763–772, 2021.
- [23] Linwei Chen, Ying Fu, Lin Gu, Chenggang Yan, Tatsuya Harada, and Gao Huang. Frequency-aware feature fusion for dense image prediction. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12):10763–10780, 2024.
- [24] Dongfang Wang, Tao Zhou, Yizhe Zhang, Shangbing Gao, and Jian Yang. Frequency-aware interaction network for ultrasound image segmentation. *IEEE Transactions on Circuits and Systems for Video Technology*, 35(7):7020–7032, 2025.
- [25] Ke Zhou, Tianxiang Chen, Jiayuan Huang, Dongmei Fu, and Chuanjiang Qi. Shape-aware medical image segmentation via frequency domain partitioning. *Engineering Applications of Artificial Intelligence*, 172:114324, 2026.
- [26] Xingguo Zhu, Wei Wang, Chen Zhang, and Haifeng Wang. Polyp-Mamba: A hybrid multi-frequency perception gated selection network for polyp segmentation. *Information Fusion*, 115:102759, 2025.
- [27] Nhat-Tan Bui, Dinh-Hieu Hoang, Quang-Thuc Nguyen, Minh-Triet Tran, and Ngan Le. MEGANet: Multi-scale edge-guided attention network for weak boundary polyp segmentation. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 7985–7994, 2024.
- [28] Xing-Liang Pan, Ju-Rong Ding, Xia Li, Shuo Liu, Jie Wang, Bo Hua, Guo-Zhi Tang, and Chang-Hua Zhong. MSBP-Net: A multi-scale boundary prediction network for automated polyp segmentation. *Pattern Recognition*, 170:112101, 2026.
- [29] Bo-Cheng Hu, Ge-Peng Ji, Dian Shao, and Deng-Ping Fan. PraNet-V2: Dual-supervised reverse attention for medical image segmentation. *Computational Visual Media*, 12(2):493–500, 2026.
- [30] Xiaolu Kang, Zhuoqi Ma, Kang Liu, Yunan Li, and Qiguang Miao. Modeling multi-scale uncertainty with evidence integration for reliable polyp segmentation. *Neural Networks*, 189:107553, 2025.
- [31] Debesh Jha, Pia H. Smedsrud, Michael A. Riegler, Pål Halvorsen, Thomas de Lange, Dag Johansen, and Håvard D. Johansen. Kvasir-SEG: A segmented polyp dataset. In *MultiMedia Modeling*, pages 451–462. Springer, 2020.
- [32] Jorge Bernal, F. Javier Sánchez, Gloria Fernández-Esparrach, Debora Gil, Cristina Rodríguez, and Fernando Vilariño. WM-DOVA maps for accurate polyp highlighting in colonoscopy: Validation versus saliency maps from physicians. *Computerized Medical Imaging and Graphics*, 43:99–111, 2015.

- [33] David Vázquez, Jorge Bernal, F. Javier Sánchez, Gloria Fernández-Esparrach, Antonio M. López, Adriana Romero, Michal Drozdal, and Aaron Courville. A benchmark for endoluminal scene segmentation of colonoscopy images. *Journal of Healthcare Engineering*, 2017:1–9, 2017.
- [34] Nima Tajbakhsh, Suryakanth R. Gurudu, and Jianming Liang. Automated polyp detection in colonoscopy videos using shape and context information. *IEEE Transactions on Medical Imaging*, 35(2):630–644, 2016.
- [35] Juan Silva, Aymeric Histace, Olivier Romain, Xavier Dray, and Bertrand Granado. Toward embedded detection of polyps in WCE images for early diagnosis of colorectal cancer. *International Journal of Computer Assisted Radiology and Surgery*, 9(2):283–293, 2014.
- [36] Deng-Ping Fan, Ming-Ming Cheng, Yun Liu, Tao Li, and Ali Borji. Structure-measure: A new way to evaluate foreground maps. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 4558–4567, 2017.
- [37] Ran Margolin, Lihi Zelnik-Manor, and Ayellet Tal. How to evaluate foreground maps. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 248–255, 2014.
- [38] Deng-Ping Fan, Cheng Gong, Yang Cao, Bo Ren, Ming-Ming Cheng, and Ali Borji. Enhanced-alignment measure for binary foreground map evaluation. In *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence*, pages 698–704, 2018.
- [39] Federico Perazzi, Philipp Krähenbühl, Yael Pritch, and Alexander Hornung. Saliency filters: Contrast based filtering for salient region detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 733–740, 2012.